

Computer Systems

Lecture 7

Computer Architecture

Dr José Cano, Dr Lito Michala
School of Computing Science
University of Glasgow
Spring 2020

Outline

- Some history: early computing machines
- Computer Architecture
- Instructions
- Memory

How computers developed

- Adding machines
- Machines that could add, subtract, multiply, divide
- Machines that could do fixed sequences of arithmetic
- Machines that could look at the results of arithmetic and then decide what to do next


Machine arithmetic

- Machines have been used to help with arithmetic for along time
 - Some assist the human (abacus)
 - Others perform arithmetic mechanically (Pascal's adder)
- Arithmetic with gears (19th century technology)
 - Video showing how carry propagation can be done with gears
<http://www.youtube.com/watch?v=YXMuJco8onQ>
 - Video of Pascal calculator
<http://www.youtube.com/watch?v=3h71HAJWnVU>
- Binary arithmetic with marbles (just for fun!)
 - <http://www.youtube.com/watch?v=GcDshWmhF4A>
 - <http://www.youtube.com/watch?v=md0TISjlags>

Three key ideas


- Make a machine that performs a **fixed sequence** of arithmetic operations
- Provide a way to **change that sequence**
 - Configure the machine for a specific problem by defining the sequence of operations to perform
 - **Jacquard loom**: used punched cards
 - **Babbage's Difference Engine**: used fixed sequences of arithmetic operations to calculate functions for nautical tables
- Make it possible to compare two numbers, and then **decide what to do next**
 - **Babbage's Analytical Engine**: first general “Turing-complete” computer
 - Instead of computing a fixed sequence of operations, there is something like a conditional **if $x < y$ then ... else ...**

Charles Babbage (1791-1871)



“Inventor of the computer”

Babbage's Difference Engine

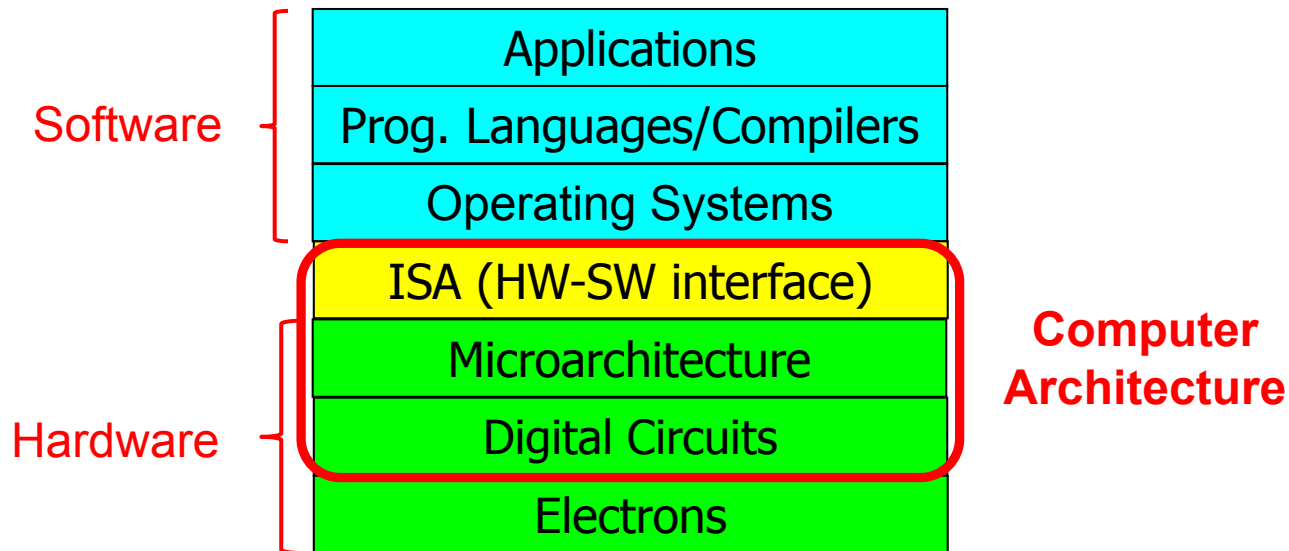


Outline

- Some history: early computing machines
- Computer Architecture
- Instructions
- Memory

Computer Architecture

- CA defines the **structure** and the **machine language** of a computer



- Understanding the principles of CA is central to computer science
- **Practical benefit:** understanding machine language will give you a deeper understanding of programming languages

The Language of Machines

- There are many programming languages
 - Python, Java, C, etc
- Each machine has one fixed **machine language**
- Why are there many programming languages for a given machine?
 - Each PL is **translated** to machine language by software called a **compiler**
- What is a computer?
 - A digital circuit (a piece of hardware, a machine) that executes programs

Machine language

- A fixed digital circuit can execute one fixed **machine language**
- Examples
 - X86 (Intel, AMD)
 - ARM
 - MIPS
 - Sparc
 - ... and many more
- The details of different machine languages are quite different
- But we will focus on principles common to all of them

What is machine language like?

- Very different from Python, Java, C, ...
- The designer of a machine language has to “look both up and down”
 - **Looking up:** the machine language must be **powerful enough** to provide the foundation for operating systems and programming languages
 - **Looking down:** the machine language must be **simple enough** so that a digital circuit can execute it (example: RTM circuit)
- Machine languages are also designed to make high performance possible
- But they are not intended to make programming easy!
- Today we'll see **assembly language**: a simplified notation of **machine language**

Outline

- Some history: early computing machines
- Computer Architecture
- **Instructions**
- Memory

Instructions

- A machine language provides instructions
- Analogous to statements in a programming language, with some differences
 - Statements can be complex: $x := 2 * (a + b/c)$
 - Instructions are simple: $R2 := R1 + R3$
 - Each instruction performs just one operation

Sigma16

- In this course we will study an architecture called Sigma16
- Designed to support several research projects at the University of Glasgow
- It's a research machine, not a commercial product
- There is a complete design, including a full digital circuit
- It hasn't been manufactured, but there are two implementations in software
 - An emulator, which you will use in this course
 - A simulator for the circuit

Why use Sigma16?

- Our focus is on **ideas** and **principles**
- Sigma16 illustrates all the main ideas, but avoids unnecessary complexity
- Example
 - Sigma16 has just one word size (16 bits) while commercial machines provide many
 - Most commercial computers have **backward compatibility** with previous versions, leading to great complexity
- Legacy architectures use an approach called **complex instruction set (CISC)**
- Simpler **reduced instruction set (RISC)** gives better performance
 - Sigma16 uses this

Structure of a computer

- All computers have several main subsystems
 - The **register file** is a set or array of “registers”
 - The RTM circuit has 4, sigma16 has 16 (R0, R1, R2, ..., R15)
 - A register is a circuit that can remember a 16-bit word
 - The **ALU** (arithmetic and logic unit) is a circuit that can do arithmetic, such as addition, subtraction, comparison, and some other operations
 - The **memory** can hold a large number of words: similar to the register file, but significantly slower and much larger
 - The **Input/Output** can transfer data from the outside world to/from the memory

Register file

- There are 16 registers
- Each register holds a 16-bit word
- We'll write the words using hexadecimal

R0	0000
R1	fffe
R2	13c4
⋮	⋮
R14	03c8
R15	0020

What are the registers actually?

- Recall the **register transfer machine circuit**
- It contains four registers, each holding 4 bits
- Sigma16 is just the same, but with 16 registers and each holds 16 bits
- Each 16-bit register is 16 copies of the **reg1** circuit
- Why program with registers, not variables like sum, count, x, etc?
 - In machine language we are programming directly with the hardware in the computer

The RTM instructions

- The RTM circuit can execute two instructions

$R2 := R1 + R0$; add two registers and load result

$R1 := 8$; load a constant

- We'll begin with the corresponding Sigma16 instructions

The add instruction

- Think of the registers as variables

- Examples

add R5,R2,R3 ; means $R5 := R2 + R3$

add R12,R1,R7 ; means $R12 := R1 + R7$

- General form:

- add dest,op1,op2 where dest, op1, op2 are registers
- The two operands are added, the result is placed in the destination
- Meaning: $dest := op1 + op2$

- Everything after a semicolon ; is a comment

Registers can hold variables

- We often think of a variable as a **box that can hold a number**
 - A register can hold a variable!
- An add instruction (or sub, mul, div) is like an assignment statement
 - add R2,R8,R2 means $R2 := R8 + R2$
 - Evaluate the right hand side $R8 + R2$
 - The operands (R8, R2) are not changed
 - Overwrite the left hand side R2 (destination) with the result
 - The old value of the destination is destroyed
 - It is **not a mathematical equation**, it is a **command** to do an operation and put the result into a register, overwriting the previous contents
- Assignment is often written $R2 := R8 + R2$
 - The $:=$ operator means assign, and does not mean equals

Notation and terminology

- Why write a notation like add R5,R2,R3 instead of $R5 := R2 + R3$?
 - It's actually more consistent because every instruction will be written in this form: a keyword for the **operation**, followed by the **operands**
 - The notation is related closely to the way instructions are represented in memory, which we'll see later

A simple program

- **Problem:** given three integers in R1, R2, R3
- **Goal:** calculate the sum $R1+R2+R3$ and put it in R4

- **Solution**

add R4,R1,R2 ; R4 := R1+R2 (this is a comment)

add R4,R4,R3 ; R4 := (R1+R2) + R3

More arithmetic instructions

- There are instructions for the basic arithmetic operations

add R4,R11,R0 ; R4 := R11 + R0

sub R8,R2,R5 ; R8 := R2 - R5

mul R10,R1,R2 ; R10 := R1 * R2

div R7,R2,R12 ; R7 := R2 / R12

- Every arithmetic operation takes its operands from registers, and loads the result into a register

Example

- Suppose we have variables a, b, c, d
- R1=a, R2=b, R3=c, R4=d
- We wish to compute $R5 = (a+b) * (c-d)$

add R6,R1,R2 ; R6 := a + b

sub R7,R3,R4 ; R7 := c - d

mul R5,R6,R7 ; R5 := (a+b) * (c-d)

- Good comments make the code easier to read!

General form of arithmetic instruction

- General form: `op d,a,b`

`op`: operation (+, −, ×, ÷)

`d`: destination register (where the result goes)

`a`: first operand register

`b`: second operand register

- **Meaning:** $R_d := R_a \text{ (op) } R_b$

- **Example:** `add R5,R2,R12` ; $R_5 := R_2 + R_{12}$

Register R0 and R15 are special!

- You should not use R0 or R15 to hold ordinary variables!
- **R0** always contains 0
 - Any time you need the number 0, it's available in R0
 - You cannot change the value of R0
 - `add R0,R2,R3` ; does nothing (R0 will not change)
 - `add R5,R2,R3` ; you can change all other registers
 - It is **legal** to use R0 as the destination, but it will still be 0 after you do it!
- **R15** holds status information
 - Some instructions place additional information in R15 (is the result negative? was there an overflow?)
 - Therefore the information in R15 is transient
 - R15 is for temporary information (not a safe place to keep long-term data)

Limitation of register file: it's small

- The register file is used to perform calculations
- In computing something like $x := (2*a + 3*b) / (x-1)$
 - All the arithmetic will be done using the register file
- But it has a big limitation
 - There are only 16 registers
 - And most programs need more than 16 variables!
- **Solution:** the **memory** is large and can hold far more data than the register file

Outline

- Some history: early computing machines
- Computer Architecture
- Instructions
- **Memory**

Memory

- The memory is similar to the register file
 - It is a large collection of words
- A variable name (x, sum, count) refers to a word in memory
- Some differences between memory and register file
 - The memory is much larger: 65,536 locations (the register file has only 16)
 - The memory cannot do arithmetic
- So our strategy in programming
 - Keep data permanently in memory
 - When you need to do arithmetic, copy a variable from memory to a register
 - When finished, copy the result from a register back to memory

Registers and memory

- Register file
 - 16 registers
 - Can do arithmetic, but too small to hold all your variables
 - Each register holds a 16-bit word
 - Names are R0, R1, R2, ..., R15
 - You can do arithmetic on data in the registers
 - Use registers to hold data temporarily that you're doing arithmetic on
- Memory
 - 65,536 memory locations
 - Each memory location holds a 16-bit word
 - Each memory location has an address 0, 1, 2, ..., 65,535
 - The machine cannot do arithmetic on a memory location
 - Use memory locations to store program variables permanently and the program itself

Copying a word between memory and register

- There are two instructions for accessing the memory
 - **load** copies a variable from memory to a register
 - **load R2,x[R0]** copies the variable x from memory to register R2
 - $R2 := x$
 - R2 is changed; x is unchanged
 - **store** copies a variable from a register to memory
 - **store R3,y[R0]** copies the word in register R3 to variable y in memory
 - $y := R3$
 - y is changed; R3 is unchanged
- Notice that we write **[R0]** after a variable name
 - Later we'll see the reason

An assignment statement in machine language

- $x := a+b+c$

```
load R1,a[R0]      ; R1 := a
load R2,b[R0]      ; R2 := b
add R3,R1,R2       ; R3 := a+b
load R4,c[R0]      ; R4 := c
add R5,R3,R4       ; R5 := (a+b) + c
store R5,x[R0]     ; x := a+b+c
```

- Use **load** to copy variables from memory to registers
- Do **arithmetic** with add, sub, mul, div
- Use **store** to copy result back to memory

Why do we have registers and memory?

- The programmer has to keep track of which variables are currently in registers
- You have to use load and store instructions to copy data between the registers and memory
- Wouldn't it be easier just to get rid of the distinction between registers and memory?
 - Do all the arithmetic on memory
- Short answer
 - Yes, it's possible to design a computer that way
 - But it makes the computer very much slower!
 - A computer without load and store instructions (where you do arithmetic on memory locations) would run between 100 and 1,000 times slower

Constants: the lea instruction

- The RTM has an instruction that loads a constant into a register
- Use the **lea** instruction
 - **lea R2,57[R0]** loads the constant 57 into R2 ; R2 := 57
 - Actually lea does more (later we'll see some advanced applications)
 - General form: **lea Rd,const[R0]**
 - Write [R0] after the constant; we'll see the reason for this later on
- **Example:** R3 := R1 + 39*R2

```
lea R4,39[R0]      ; R4 := 39
mul R3,R4,R2       ; R3 := 39 * R2
add R3,R1,R3       ; R3 := R1 + (39*R2)
```

Stopping the program

- The last instruction of every program should be

```
trap R0,R0,R0    ; halt
```

- This tells the computer to halt, it stops execution of the program

Defining variables

- To define variables x, y, z and give them initial values

x data 34 ; x is a variable with initial value 34

y data 9 ; y is initially 9

z data 0 ; z is initially 0

abc data \$02c6 ; specify initial value as hex

- The data statements should come after all the instructions in the program

Reason: the booter

A complete example program

- ; Program Add
- ; A minimal program that adds two integer variables
- ; Execution starts at location 0, where the first instruction will be
- ; placed when the program is executed


```
load R1,x[R0]      ; R1 := x
load R2,y[R0]      ; R2 := y
add R3,R1,R2       ; R3 := x + y
store R3,z[R0]     ; z := x + y
trap R0,R0,R0      ; terminate
```

- ; Static variables are placed in memory after the program

```
x data 23
```

```
y data 14
```

```
z data 99
```



<https://xkcd.com/1437/>