

Computer Systems

Lecture 12

Procedures and the Call Stack


Dr Lito Michala, Dr José Cano
School of Computing Science
University of Glasgow
Spring 2020

Outline

- Procedures
- Call and return
- Parameter passage
- Procedure calls another procedure
- Saving state
- Stack

Pointers quick-review

- A pointer is an address
 - much like your home address on google maps
- A value is what the pointer points at
 - much like you, in your home, on google maps
- Pointers and values are both stored in memory



Procedures: reusable code

- Often there is a sequence of instructions that comes up again and again
 - For example: sqrt (square root)
 - It takes a lot of instructions to calculate a square root
 - An application program may need a square root in many different places
- We don't want to keep repeating the code
 - It's tedious
 - It wastes space (all those instructions require memory!)
- The aim: write it once and reuse the same instructions many times

Procedure

- Write the code **one time**: the block of code is called a procedure (or subroutine, or function)
- Put the instructions by themselves somewhere, not in the main flow of instructions
- Give the block of code a label (e.g. work) that describes what it does
- Every time you need to perform this computation, **call it**: *goto work*
- When it finishes, **the procedure needs to return**
 - Go back to **the instruction after** the one that jumped to it

Outline

- Procedures
- Call and return
- Parameter passage
- Procedure calls another procedure
- Saving state
- Stack

Call and return

- One idea is just to use jump instructions for both call and return
- But that isn't actually sufficient
 - Let's look in more detail at what happens

Returning to the instruction after the call

- Suppose a procedure named **dowork** is used in several places
- Each call jumps to the same place (the address of the first instruction of the procedure)
- But the calls **come from different places**
- Therefore the procedure must finish by **returning to different places**

Calling and returning

- Here is a main program that calls a procedure “dowork” several times
 - It takes the value in R1 and doubles it, and the main program would use the result but we ignore that here

```
graph TD
    L1[load R1,x[R0] ; R1 = x]
    L2["goto" dowork[R0] ; call the procedure "dowork"]
    L3[load R1,y[R0] ; R1 = y]
    L4["goto" dowork[R0] ; call dowork]
    L5[sub R5,R6,R7 ; R5 = R6-R7]
    L6[add R1,R1,R1 ; R1 = R1+R1]
    L7["return" ; return go back to the caller]
    
    L2 -- red arrow --> L6
    L4 -- red arrow --> L6
    L7 -- green arrow --> L2
```

load R1,x[R0] ; R1 = x

“goto” dowork[R0] ; call the procedure “dowork”

load R1,y[R0] ; R1 = y

“goto” dowork[R0] ; call dowork

sub R5,R6,R7 ; R5 = R6-R7

add R1,R1,R1 ; R1 = R1+R1

“return” ; return (go back to the caller)

- The return must go to different points in the program: How can it do that?

The jump-and-link instruction: jal

- When the main program calls the subroutine, it needs to remember where the call came from
- This is the purpose of the jal instruction (jump and link)
- jal R5,dowork[R0]
 - A pointer to the next instruction after the jal (the return address) is loaded into the destination register (e.g. R5)
 - Then the machine jumps to the effective address

Jumping

- All jump instructions (jump, jal, jumplt, etc) refer to **effective addresses**

- jump **loop**[R0]

goto loop

- jump **0**[R14]

goto instruction whose address is in R14

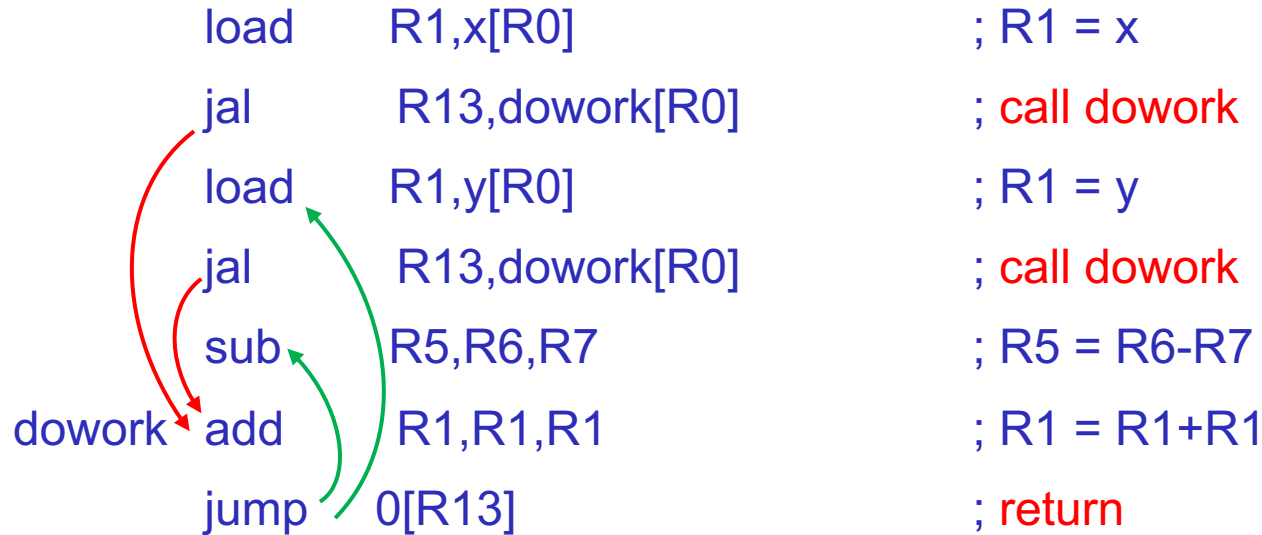
- jump **const**[R2]

goto instruction whose address is const+R2

Implementing call and return

- To call a procedure dowork: `jal R13,dowork[R0]`
 - The address of the instruction after the jal is placed in R13
 - The program jumps to the effective address, and the procedure starts executing
- To return when the procedure has finished: `jump 0[R13]`
 - The effective address is $0 + \text{the address of the instruction after the jal}$
 - The program jumps there and the main program resumes

Calling with jal and returning with jump



- **call:** jal puts a pointer to the next instruction into R13
- **return:** follow the pointer in R13

Outline

- Procedures
- Call and return
- **Parameter passage**
- Procedure calls another procedure
- Saving state
- Stack

Parameter passage

- There are several different conventions for passing argument to the function, and passing the result back
- What is important is that the caller and the procedure agree on how information is passed between them
- If there is a small number of arguments, the caller may put them in registers before calling the procedure
- If there are many arguments, the caller builds an array or vector (sequence of adjacent memory locations), puts the arguments into the vector
- and passes the address of the vector in a register (typically R1)
- A simple convention: **the argument and result are passed in R1**

Functions

- A **function** is a procedure that
 - Receives a parameter (a word of data) from the caller
 - Calculates a result
 - Passes the result back to the caller when it returns
- A **pure function** is a function that doesn't do anything else
 - It doesn't change any global variables, or do any input/output

Example: Passing argument and result in R1

; Main program

load R1,x[R0] ; arg = x

jal R13,work[R0] ; result = work (x)

...

load R1,y[R0] ; arg = y

jal R13,work[R0] ; result = work (y)

...

; Function work (x) = 1 + 7*x

work lea R2,7[R0] ; R7 = 2

lea R3,1[R0] ; R3 = 1

mul R1,R1,R2 ; result = arg * 7

add R1,R3,R1 ; result = 1 + 7*arg

jump 0[R13] ; return

Outline

- Procedures
- Call and return
- Parameter passage
- Procedure calls another procedure
- Saving state
- Stack

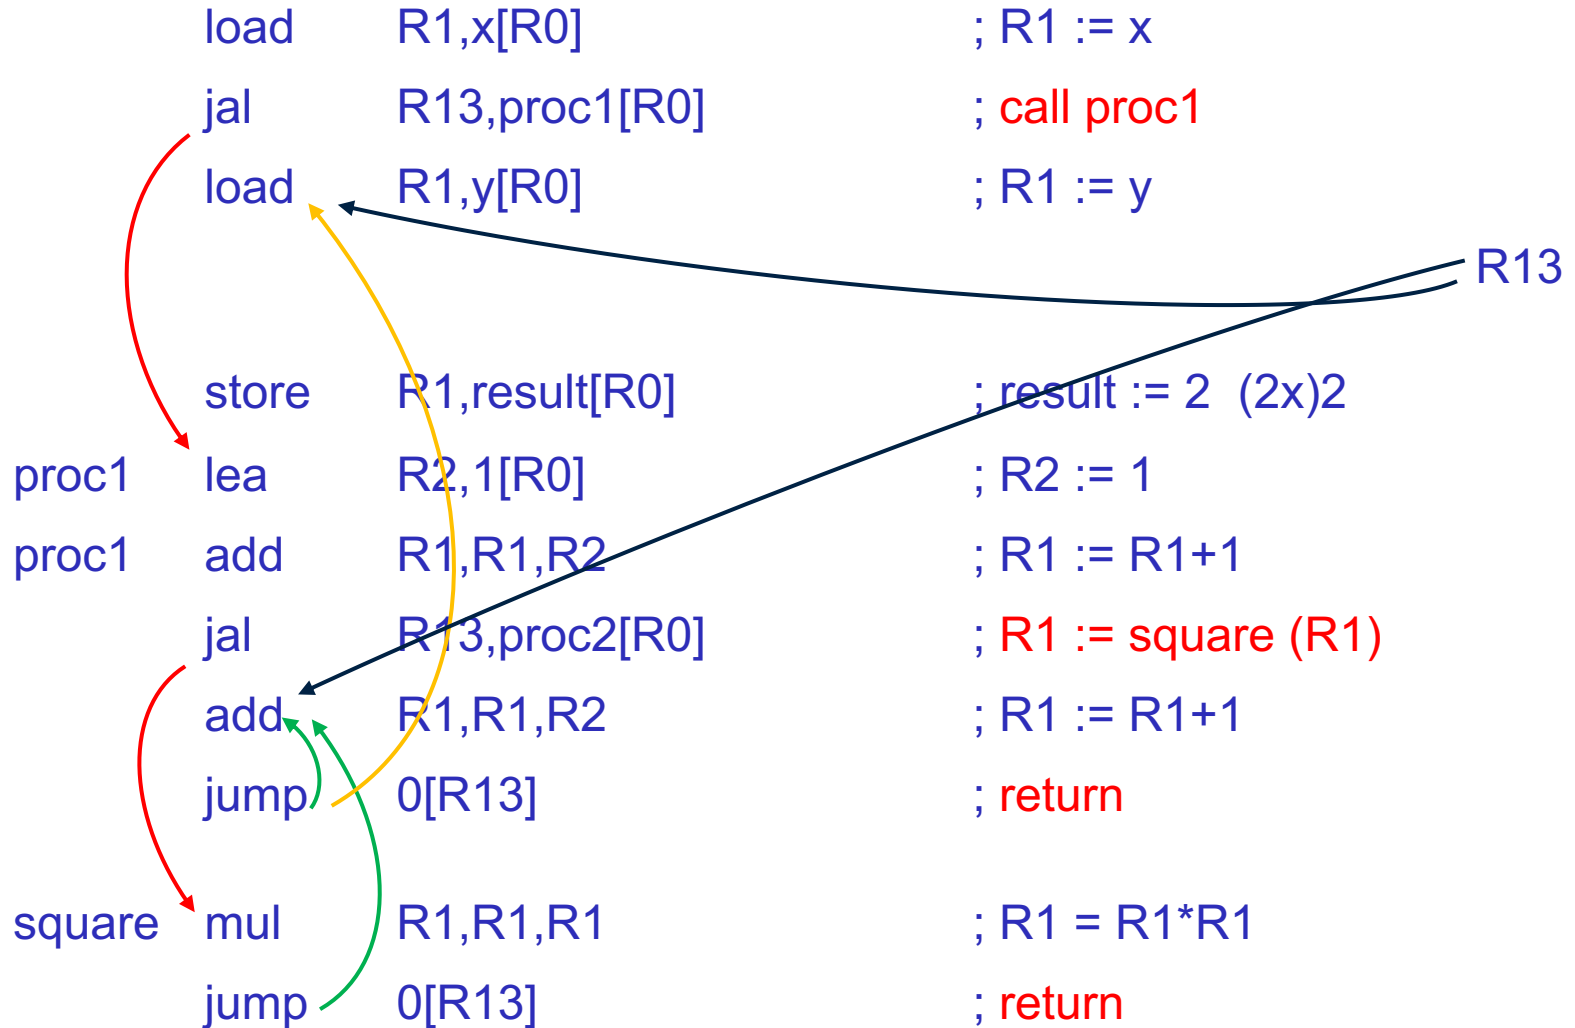
What if a procedure calls another procedure?

- The simplest kind of procedure
 - Call it with `jal R13,procname[R0]`
 - It returns by executing `jump 0[R13]`

Limitations of basic call

- If the procedure modifies any registers, it may destroy data belonging to the caller
- If the procedure calls another procedure, it can't use R13 again
 - Each procedure would need a dedicated register for its return address, limiting the program to a small number of procedures
- The basic call mechanism doesn't allow a procedure to call itself
 - This is called **recursion**

R13 overwritten: proc1 returns to the wrong place!



Outline

- Procedures
- Call and return
- Parameter passage
- Procedure calls another procedure
- Saving state
- Stack

Saving state

- Calling a procedure creates new information
 - The return address
 - Whatever values the procedure loads into the registers
- But this new information could overwrite essential information belonging to the caller
- We need to **save the caller's state** so the procedure won't destroy it

The wrong way to save state

- Suppose we just have a variable saveRetAdr
- Store R13 into it in the procedure, load that when we return
- Now it's ok for proc1 to call proc2
- But if proc2 calls proc3 we are back to the same problem: it doesn't work!
- The solution: a **stack**

Saving registers

- Most procedures need to use several registers
- It's nearly impossible to do anything without using some registers!
- The first thing a procedure should do is to save the registers it will use by copying them into memory (with store instructions)
- The last thing it should do before returning is to restore the registers by copying their values back from memory (with load instructions)

Where can the registers be saved?

- It won't work to copy data from some of the registers to other registers!
- It's essential to save the data into memory
- Two approaches
 - Allocate fixed variables in memory to save the registers into (simple but doesn't allow recursion)
 - Maintain a **stack** in memory, and **push** the data onto the stack (this is the best approach and is used by most programming languages)

Who saves the state: the caller or the procedure?

- Two approaches
- Caller saves (used occasionally)
 - Before calling a procedure, the caller saves the registers, so all its essential data is in memory
 - After the procedure returns, the caller does whatever loads are needed
- Callee saves (usually the preferred solution)
 - The caller keeps data in registers, and assumes that the procedure won't disturb it
 - The first thing the procedure does is to save the registers it needs to use into memory
 - Just before returning, the procedure restores the registers by loading the data from memory

Outline

- Procedures
- Call and return
- Parameter passage
- Procedure calls another procedure
- Saving state
- Stack


Stack of return addresses

- To allow a large number of procedures, we can't dedicate a specific register to each one for its return address
- Therefore we
 - Always use the same register for the return address in a jal instruction (we will use R13)
 - The first thing a procedure does is to store its return address into memory
 - The last thing the procedure does is to load its return address and jump to it
 - The return addresses are pushed onto a stack, rather than being stored at a fixed address


Stacks

- A **stack** is a container
- Initially it is empty
- You can **push** a value onto the stack: this is now sitting on the **top** of the stack
- You can **pop** the stack: this removes the **most recently pushed** value and returns it
- A stack allows access only to the top value; you cannot access anything below the top
- We can save procedure return addresses on a stack because return always needs the most recently saved return address


Initially the stack is empty




Call procedure, push return address a




Call another procedure, push return address b



Return: pop produces return address b



Call some procedure, push return address c



Call a procedure, push return address d

d
c
a

The call stack

- Central technique for
 - Preserving data during a procedure call
 - Holding most of your variables
- It goes by several names (these are all the same thing)
 - call stack
 - execution stack
 - “The stack”
- It's important!
 - Most programming languages use it
 - Computers are designed to support it
 - Often referred to (Stack Overflow web site, etc.)

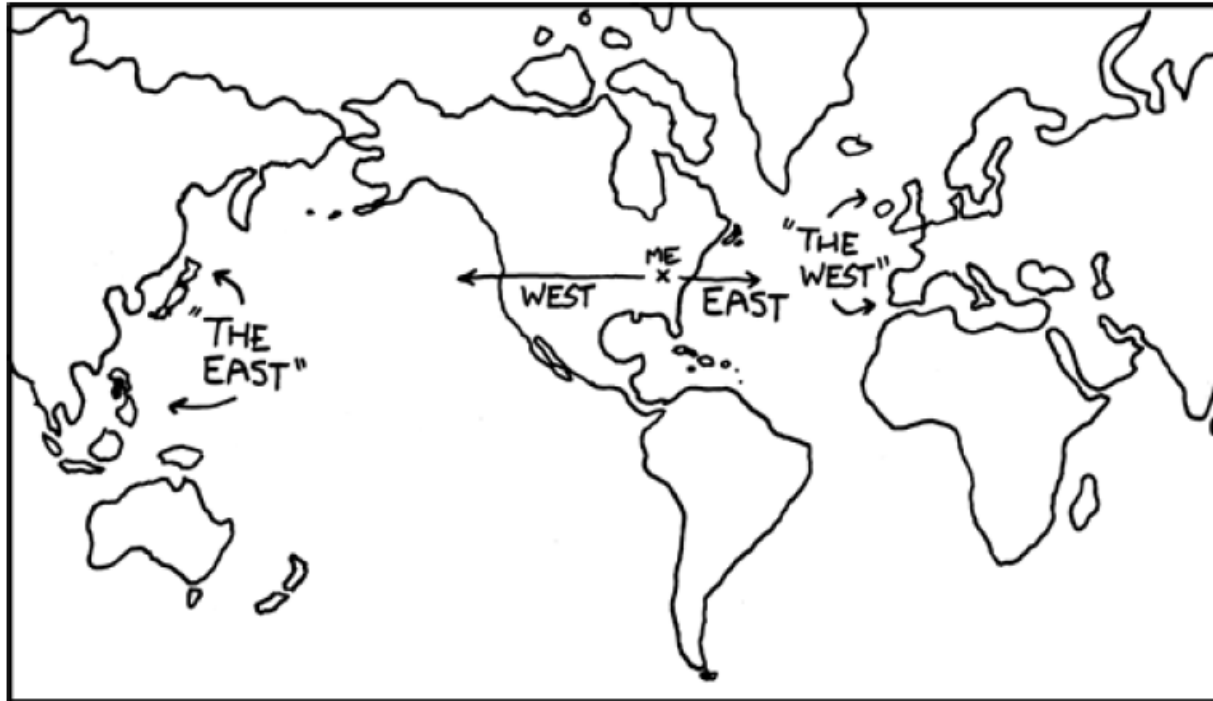
Stack frames

- There is a **call stack** or **execution stack** that maintains complete information about all procedure calls and returns
- Every “activation” of a procedure pushes a **stack frame**
- When the procedure returns, its stack frame is popped (removed) from the stack
- R14 contains the address of the current (top) stack frame
- The stack frame contains
 - A pointer to the previous stack frame (this is required to make the pop work)
 - The return address (saved value of R13)
 - The saved registers (so the procedure can use the registers without destroying information)
 - Local variables (so the procedure can have some memory of its own to use)

Implementing the call stack

- Dedicate R14 to the **stack pointer**
- This is a programming convention, not a hardware feature
- When the program is started, R14 will be set to point to an empty stack
- When a procedure is called, the saved state will be pushed onto the stack: store a word at 0[R14] and add 1
- When a procedure returns, it pops the stack and restores the state: subtract 1, load from 0[R14]
- The program should never modify R14 apart from the push and pop

terminology



THIS ALWAYS BUGGED ME.

<https://xkcd.com/503/>