

Computer Systems

Lecture 16

Nested Conditionals

Dr Lito Michala, Dr José Cano
School of Computing Science
University of Glasgow
Spring 2020

Outline

- Blocks
 - Syntax for blocks
 - Enter at beginning, exit at end
- Nested conditionals
- Case and jump tables
- Real world incident: the SSL Bug

Blocks

- You'll need to learn many programming languages, eventually
 - There are concepts that appear in most languages
 - It's useful to focus on the general concept
- One important concept is a **block**
 - A block or **compound statement** is a single statement that contains several statements
 - The purpose is to let you have a group of statements in a loop, or controlled by a conditional
- The **syntax** is the detailed punctuation used to indicate a block, and this varies in different languages
 - There are lots of small syntax differences between languages
 - Some languages use `:=` to mean assign, `=` to mean equals
 - Some languages use `=` to mean assign, `==` to mean equals

Python syntax style for blocks: layout

- The layout (the indentation) determines what is inside the if statement

```
a = 1
```

```
if x < y :
```

```
    b = 2
```

```
    c = 3
```

```
d = 4
```

- If you change the indentation, you change the meaning of the program
- In Python, you write : to mean then, and else to mean else

Algol/Pascal style for blocks: begin-end

```
a := 1;  
if x < y  
    then begin b := 2;  
            c := 3  
        end  
    else begin d := 4;  
            e := 5  
        end;  
f := 6;
```

- You write then to mean then, and else to mean else
- Statements must be separated by semicolon ;

C style for blocks: braces

```
a = 1;  
if (x < y) b = 2;  
else  
    { d = 4;  
      e = 5; }  
f = 6;
```

- In C you don't write then at all, but this means the condition $x < y$ must be enclosed in parentheses so the compiler can tell where the condition ends and the then-statement begins

Block structured style for blocks: matching keywords

```
a := 1;  
if x < y  
    then b := 2;  
        while i < n do  
            sum := sum + x[i];  
            i := i + 1  
        endwhile  
    else d := 4  
        e := 5  
endif;  
f := 6;
```

- This style introduces a lot of keywords (endif, endwhile, endfor, endrepeat) but you don't need braces around a block
 - It makes code more readable and enables the compiler to produce better error messages

Enter at beginning, exit at end

- A common programming style is to require
 - Each block enters only at the beginning of the block
 - Each block exits only at the end of the block
- This style is helpful in some programming languages, but in some languages it sometimes makes code less readable
- In assembly language, and for compilation patterns, it is necessary to follow this style
- In high level languages this style is sometimes helpful, but not always

Single entrance/exit for compilation patterns

- It is straightforward to translate high level control constructs using the compilation patterns
- These require that the blocks of code always start at the beginning and finish at the end
- It's bad to jump into the middle of a block, or exit out of the middle because
 - The compilation patterns won't work correctly
 - You'll have to duplicate a lot of code
 - Example: returning from a procedure requires restoring the registers, resetting the stack pointer, loading the return address
 - That code should not be duplicated in several places in a procedure

Systematic approach to programming?

- Start by understanding what your program should do
- Express the algorithm using high level language notation (and it's ok to mix in some English too)
- Translate the high level to the low level
 - Assignment statements: $x := \text{expression}$
 - I/O statements: Write string
 - goto L
 - if boolean then goto L
- Translate the low level to assembly language
- Retain the high and low level code as comments
- Do hand execution at every level

Why use this systematic approach to programming?

- It enables you to write correct code at the outset, and minimize debugging
- If there is a bug, it helps you to catch it early (e.g. in translation to goto form)
- If there's a bug in an instruction, the comments enable you to find it quickly
 - A common error is to use a wrong register number: `add R9,R3,R4`
 - Poor comments don't help: `add R9,R3,R4 ; R9 := R3+R4`
 - Good comments help a lot: `add R9,R3,R4 ; x := alpha + (a[i]*b[i])`
 - Look at the register usage comments (oops, x is in R8, not R9, now I know how to x it and I don't have to read the entire program)
- Professional software needs to be maintained
 - The comments make the software easier to read and more valuable

Outline

- Blocks
 - Syntax for blocks
 - Enter at beginning, exit at end
- Nested conditionals
- Case and jump tables
- Real world incident: the SSL Bug

Nested if-then-else

- Conditional statements can be nested deeply

```
if b1
  then S1
    if b2
      then S3
      else S4
    S5
  else S6
    if b3
      then S7
    S8
```

Special case of nested if-then-else

- Often the nesting isn't random, but has this specific structure

```
if b1
  then S1
  else if b2
    then S2
    else if b3
      then S3
      else if b4
        then S4
        else if b5
          then S5
          else S6
```

- (Some languages require punctuation to avoid ambiguity)

Another way to write it

- To avoid running o the right side of the window, it's usually indented like this

```
if b1
    then S1
else if b2
    then S2
else if b3
    then S3
else if b4
    then S4
else if b5
    then S5
else S6
```

Some programming languages have elsif or elif

```
if b1 then S1
  elif b2 then S2
  elif b3 then S3
  elif b4 then S4
  elif b5 then S5
  else S6
```

- It avoids ambiguity
- It signals to the compiler and to a human programmer that this specific construct is being used
- It allows good indentation layout without violating the basic principle of indentation
- Some languages have this, some don't

A common application: numeric code

- Nested if-then-else but the boolean conditions are not arbitrary
 - They are checking the value of a code

```
if code = 0
  then S1
else if code = 1
  then S2
else if code = 2
  then S3
else if code = 3
  then S4
else if code = 4
  then S5
```

Outline

- Blocks
 - Syntax for blocks
 - Enter at beginning, exit at end
- Nested conditionals
- Case and jump tables
- Real world incident: the SSL Bug

The case statement

case n of

0 -> Stmt

1 -> Stmt

2 -> Stmt

3 -> Stmt

4 -> Stmt

5 -> Stmt

else -> Stmt // handle error

- This means: execute the statement corresponding to the value of n
- Many programming languages have this; the syntax varies a lot but that isn't what's important

Example: numeric code specifies a command

; The input data is an array of records, each specifying an operation

; Command : record

; code : Int ; specify which operation to perform

; i : Int ; index into array of lists

; x : Int ; value of list element

; The meaning of a command depends on the code:

; 0 terminate the program

; 1 insert x into set[i]

; 2 delete x from set[i]

; 3 return 1 if set[i] contains x, otherwise 0

; 4 print the elements of set[i]

Selecting the command with a case statement

```
; Initialize
;   BuildHeap ()

; Execute the commands in the input data
;   finished := 0
;   while InputPtr <= InputEnd && not finished
;       CurrentCode := (*InputPtr).code
;       p := set[*InputPtr] ; linked list
;       x := (*InputPtr).x ; value to insert, delete, search

;       case CurrentCode of
;           0 : <CmdTerminate>
;           1 : <CmdInsert>
;           2 : <CmdDelete>
;           3 : <CmdSearch>
;           4 : <CmdPrint>
;           else : <>
;       InputPtr := InputPtr + sizeof(Command)
; Terminate the program
; halt
```

Finding a numeric code

- It's tedious and inefficient to go through the possible values of a numeric code in sequence
 - If you're looking up Dr Zhivago in the phone book, do you look at Arnold Aardvark, and Anne Anderson, on on and on?
 - You go straight to the end of the book
- We want to find the statement corresponding to a numeric code directly, without checking all the other values

A problem with efficiency

- The problem
 - There are many applications of case statements
 - They are executed frequently
 - The number of cases can be large (not just 4 or 5; can be hundreds)
 - The implementation of the if-then-else requires a separate compare and jump for each condition
- The solution
 - A technique called **jump tables**

Jump tables: the basic idea

- For each target statement (S1, S2, S3, etc) in the conditional, introduce a jump to it: jump S1[R0], jump S2[R0], etc
- Make an **array** “jt” of these jump instructions
 - jt[0] = jump S0[R0]
 - jt[1] = jump S1[R0]
 - jt[2] = jump S2[R0]
 - jt[3] = jump S3[R0]
 - jt[4] = jump S4[R0]
- Given the code, Jump to jt[code]
- Each element of the array is an instruction that requires two words
- So jump to &jt + 2 * code

Jump table

; Jump to operation specified by code

add R4,R4,R4

lea R5,CmdJumpTable[R0]

add R4,R5,R4

jump 0[R4]

CmdJumpTable

jump CmdTerminate[R0]

jump CmdInsert[R0]

jump CmdDelete[R0]

jump CmdSearch[R0]

jump CmdPrint[R0]

; code := 2*code

; R5 := pointer to jump table

; address of jump to operation

; jump to jump to operation

; code 0: terminate the program

; code 1: insert

; code 2: delete

; code 3: search

; code 4: print

CmdDone

load R5,InputPtr[R0]

lea R6,3[R0]

add R5,R5,R6

store R5,InputPtr[R0]

jump CommandLoop[R0]

We have to be careful!

- What if code is negative, or larger than the number of cases?
- The jump to the jump table could go anywhere!
- It might not even go to an instruction
- But whatever is in memory at the place it goes to, will be interpreted as an instruction that will be executed
- The program will go haywire
- Debugging it will be hard: the only way to catch the error is to read the code and/or to single step
- Solution: before jumping into the jump table, check to see if code is invalid (too big or too small)

Checking whether the code is invalid

CommandLoop

load R5,InputPtr[R0]	; R1 := InputPtr
load R4,0[R5]	; R4 := *InputPtr.code
; Check for invalid code	
cmp R4,R0	; compare (*InputPtr).code, 0
jumplt CmdDone[R0]	; skip invalid code (negative)
lea R6,4[R0]	; maximum valid code
cmp R4,R6	; compare code with max valid code
jumpgt CmdDone[R0]	; skip invalid code (too large)

...

CmdDone

```
load R5,InputPtr[R0]
lea R6,3[R0]
add R5,R5,R6
store R5,InputPtr[R0]
jump CommandLoop[R0]
```

Outline

- Blocks
 - Syntax for blocks
 - Enter at beginning, exit at end
- Nested conditionals
- Case and jump tables
- Real world incident: the SSL Bug

A real world incident!

- The moral of the following true story: be meticulous!
- On 21 February 2014, Apple announced a bug in core security software
- It affected iPhone, iPad, Macintosh, Safari
- The bug is serious: it skips a crucial security check in SSL
- A security update was made available quickly
 - That's why you should always apply your security updates!
- It was fixed quickly but we can learn some important lessons from it

What is SSL?

- Secure Sockets Layer
- It arranges that communication between your device and a server is encrypted
- When you see a “padlock icon” on a browser, it means the SSL protocol is being used

Source 1e: sslKeyExchange.c

- 1,969 lines of code in C
- You can see the code: click on this link, or type it in:
 - http://opensource.apple.com/source/Security/Security-55471/libsecurity_ssl/lib/sslKeyExchange.c
- The bug made it possible to break open “secure” Internet traffic
- It appears somewhere on the next three slides: **can you spot it?**

The function SSLVerifySignedServerKeyExchange (1)

```
static OSStatus
SSLVerifySignedServerKeyExchange(SSLContext *ctx, bool isRsa, SSLBuffer signedParams,
                                uint8_t *signature, UInt16 signatureLen)
{
    OSStatus      err;
    SSLBuffer      hashOut, hashCtx, clientRandom, serverRandom;
    uint8_t        hashes[SSL_SHA1_DIGEST_LEN + SSL_MD5_DIGEST_LEN];
    SSLBuffer      signedHashes;
    uint8_t        *dataToSign;
    size_t         dataToSignLen;

    signedHashes.data = 0;
    hashCtx.data = 0;

    clientRandom.data = ctx->clientRandom;
    clientRandom.length = SSL_CLIENT_SRVR_RAND_SIZE;
    serverRandom.data = ctx->serverRandom;
    serverRandom.length = SSL_CLIENT_SRVR_RAND_SIZE;
```


The function continued (2)

```
if(isRsa) {
    /* skip this if signing with DSA */
    dataToSign = hashes;
    dataToSignLen = SSL_SHA1_DIGEST_LEN + SSL_MD5_DIGEST_LEN;
    hashOut.data = hashes;
    hashOut.length = SSL_MD5_DIGEST_LEN;

    if ((err = ReadyHash(&SSLHashMD5, &hashCtx)) != 0)
        goto fail;
    if ((err = SSLHashMD5.update(&hashCtx, &clientRandom)) != 0)
        goto fail;
    if ((err = SSLHashMD5.update(&hashCtx, &serverRandom)) != 0)
        goto fail;
    if ((err = SSLHashMD5.update(&hashCtx, &signedParams)) != 0)
        goto fail;
    if ((err = SSLHashMD5.final(&hashCtx, &hashOut)) != 0)
        goto fail;
}
else {
    /* DSA, ECDSA - just use the SHA1 hash */
    dataToSign = &hashes[SSL_MD5_DIGEST_LEN];
    dataToSignLen = SSL_SHA1_DIGEST_LEN;
}
hashOut.data = hashes + SSL_MD5_DIGEST_LEN;
hashOut.length = SSL_SHA1_DIGEST_LEN;
```

The function continued (3)

```
    if ((err = SSLFreeBuffer(&hashCtx)) != 0)
        goto fail;

    if ((err = ReadyHash(&SSLHashSHA1, &hashCtx)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &clientRandom)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
        goto fail;

    err = sslRawVerify(ctx,
                      ctx->peerPubKey,
                      dataToSign,
                      dataToSignLen,
                      signature,
                      signatureLen);
                      /* plaintext */
                      /* plaintext length */

    if(err) {
        sslErrorLog("SSLDecodeSignedServerKeyExchange: sslRawVerify "
                    "returned %d\n", (int)err);
        goto fail;
    }

fail:
    SSLFreeBuffer(&signedHashes);
    SSLFreeBuffer(&hashCtx);
    return err;
}
```

A bit of C

- Syntax of the if-statement

if (condition) statement;

- Examples

if (x<y) n=n+1;

if (n==5)

{ p = 3;

q = 4;

printf ("hello p = %3d\n", p);

}

A closer look

```
if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
    goto fail;
```

```
if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
    goto fail;
    goto fail;
```

```
if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
    goto fail;
```

- The second goto fail is **unconditional** so the check in the last if statement is **never executed**

What can we learn from this?

- To be a good programmer you need to
 - Know your programming language
 - Be meticulous
 - **Read** your code carefully
 - Do thorough testing
- It was very helpful that Apple published the source code for this bug

A mystery: The Attempted Linux Trapdoor Hack of 2003

- A **trapdoor** is a flaw (bug? hack?) in software that allows a hostile user to gain abilities they should not have
 - It's like giving the key to the safe in a bank to a burglar
- The master copy of Linux was maintained on Bitkeeper
 - No longer, now git is used
- A secondary copy was on cvs
- Someone (unknown to this day) made a small edit to an obscure piece of code on the cvs version

The offending code: can you spot the trapdoor?

- A change was made to a function that a user could call to wait for an event
- System functions typically take options, and they need to check to see if the options are valid
- If not, a return code *retval* is set to a non-zero value
- So, here's the code

```
if ((options == (__WCLONE|__WALL)) && (current->uid = 0))  
    retval = -EINVAL;
```

What's going on?

- In C, == is the comparison operator = is the assignment operator (bad notation, but we're stuck with it)
- This code looks like it is checking:
 - If either the WCLONE or WALL option ag is set and the current user is root, then return an error condition
- What it **actually** says is
 - Check the irrelevant ags, **set the current user to be root**, get the irrelevant result 0 and do nothing else


```
if ((options == (__WCLONE|__WALL)) && (current->uid = 0))  
    retval = -EINVAL;  
]
```


What's the significance of root?


- Many **processes** are running in the computer
- Each has its own set of privileges
- There is one special user “root” which has all privileges: **root can do anything at all**
- The effect of this change to the source code (a change of one character! is
 - If a user program calls this obscure function with just the right set of obscure options, **it suddenly gains full control over the machine**
- Don't worry, **this faulty code never made it into the master copy of Linux**

HIGHWAY ENGINEER PRANKS:

THE INESCAPABLE CLOVERLEAF:



THE ZERO-CHOICE INTERCHANGE:



THE ROTARY SUPERCOLLIDER:

