



University
of Glasgow | School of
Computing Science


Networks & Operating Systems Essentials

Dr. Angelos Marnerides

<angelos.marnerides@glasgow.ac.uk>

School of Computing Science

Coming up next...



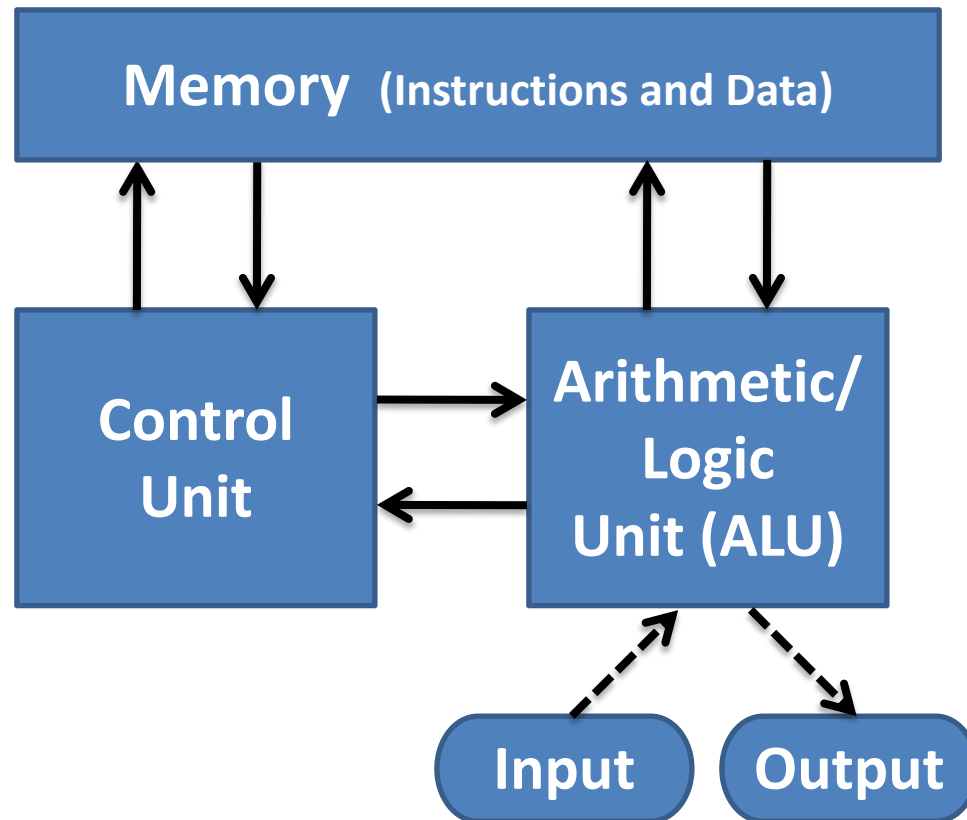
Source: <https://xkcd.com/1508/>

What is going to be covered?

- Processes
 - Process management
 - Scheduling
- Memory
 - Virtual memory
 - Page tables
- Storage
 - Block storage
 - File systems (plus case studies)

Elements of computer architecture

- Today, computer architecture is largely standardised, at a high level of abstraction, on the *von Neumann Architecture*



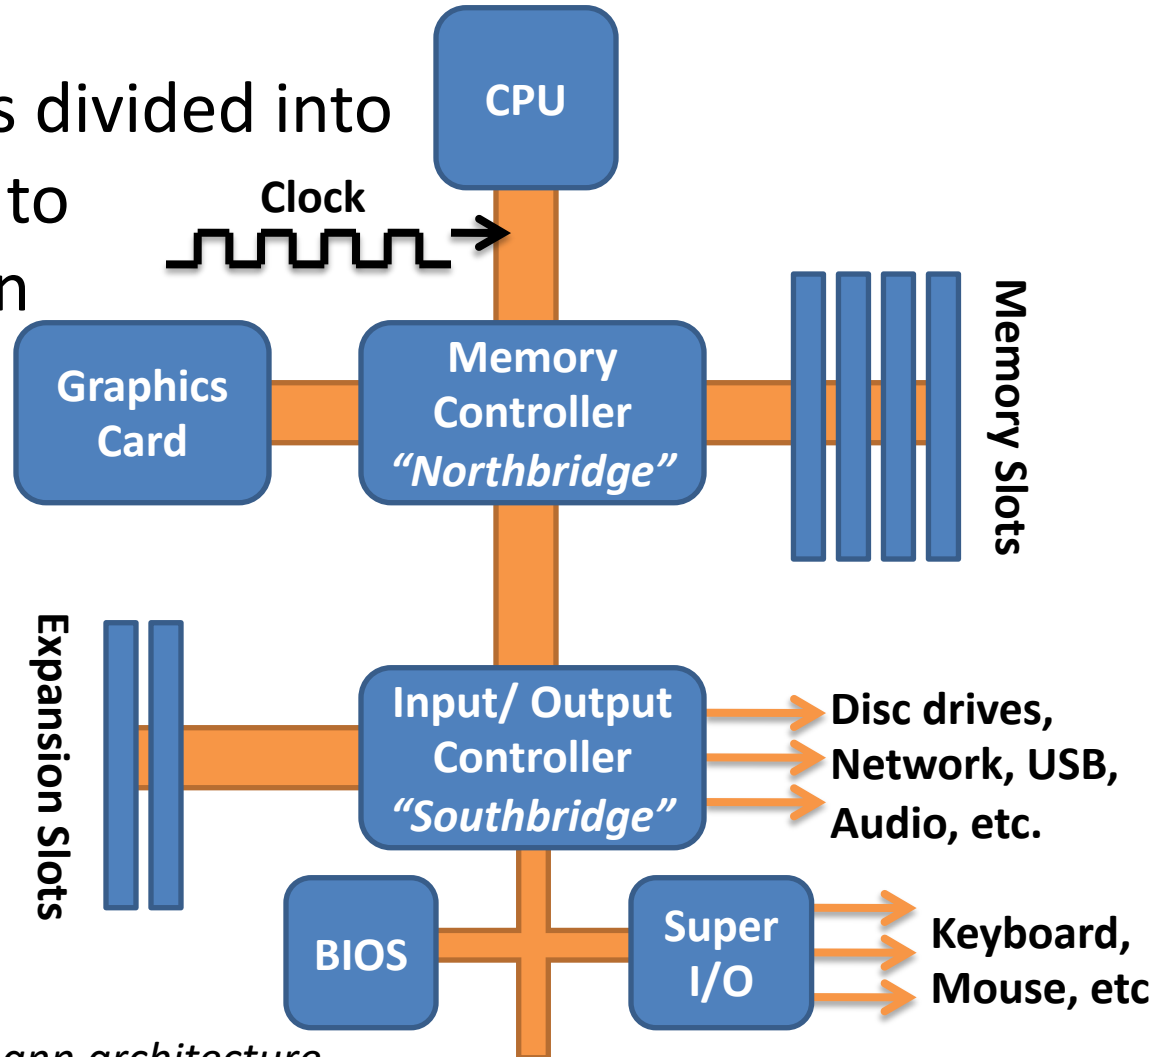
A (fairly) modern PC architecture

- The architecture is divided into *regions* according to speed of operation

Higher speed



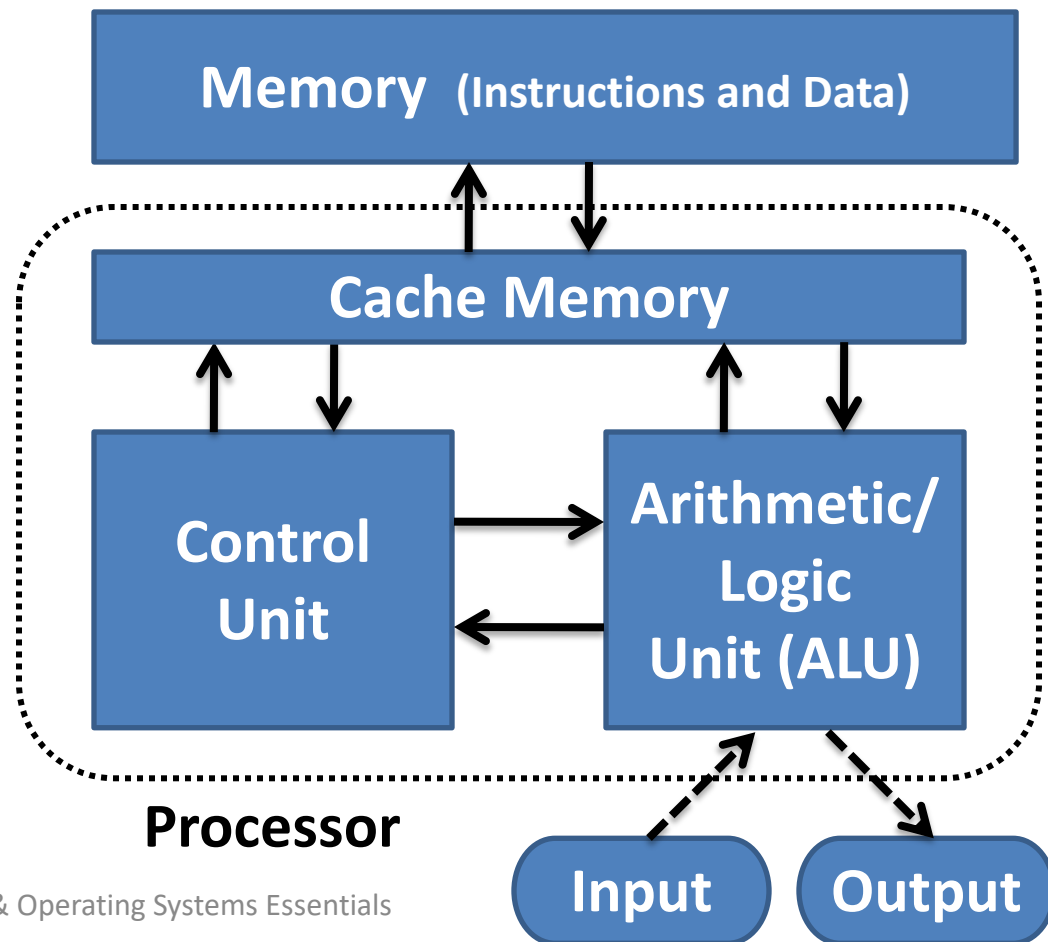
Lower speed



Still a von Neumann architecture

The processor

- Often also referred to as the Central Processing Unit (CPU)
- ALU + Control Unit; often also contains some internal high-speed *cache memory*
- Note, this is still logically the same picture as on the previous slide - still a von Neumann architecture



Instruction Set Architecture (ISA)

– The view of the processor that is seen by programs being executed by that processor

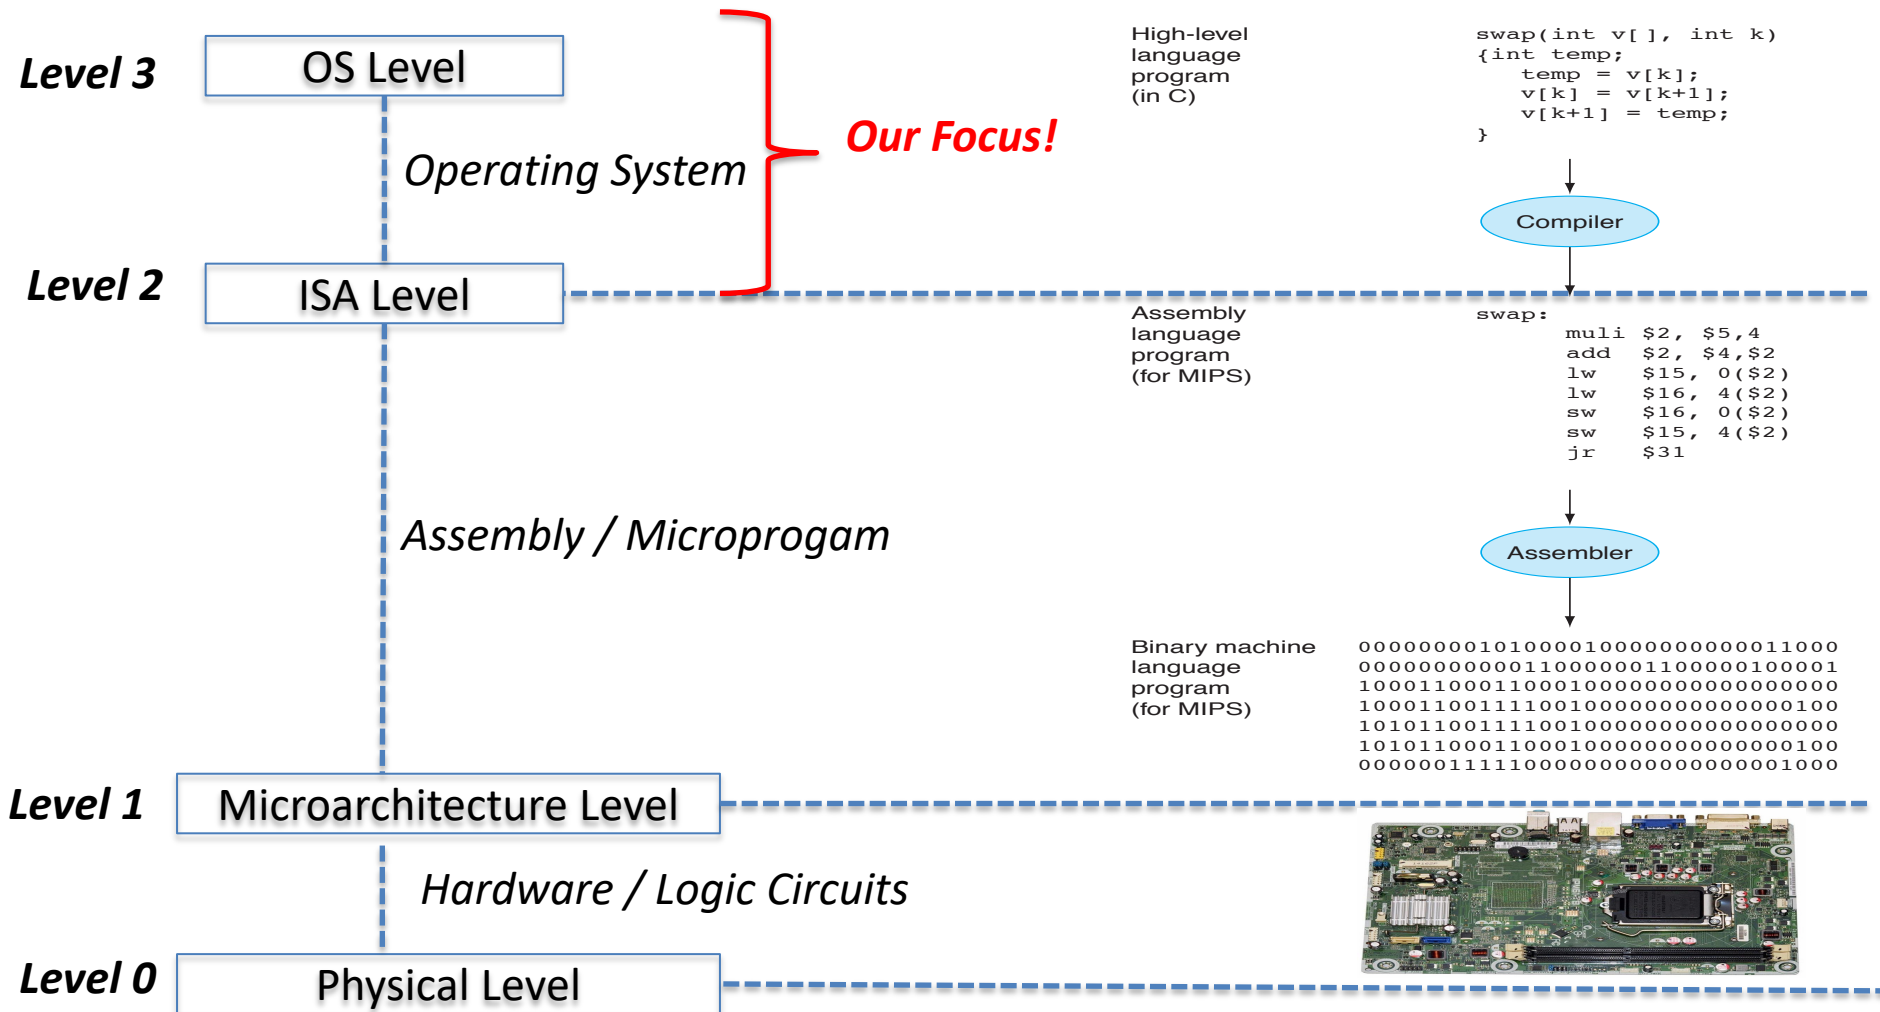
- What is the set of available instructions?
- What is the set of available registers?
- How many operands do instructions need?
- What are the sizes and types of the operands?
- How are operands accessed (e.g., stack-based ISAs don't support random access to memory – see next slide)?
- How many operands can be in registers (vs. in memory)?
- How many clock ticks does it take to execute an instruction?
- ...

```
mul $2, $5, 4
add $2, $4, $2
lw $15, 0($2)
lw $16, 4($2)
sw $16, 0($2)
sw $15, 4($2)
jr $31
```

Registers

- We've looked at instructions, now let's look more at *registers*...
- Registers are *holding areas for data being worked on inside the CPU*
 - Often, arithmetic and logic instructions are designed to work *on registers only*, not on main memory
 - This is because registers are much *faster* than main memory
 - We use the data-transfer instructions for register ↔ main memory transfers
- General purpose and special registers
 - The registers used by the arithmetic and logic instructions are called **general purpose registers**
 - In computers like MIPS we have a largish set of these—a “file” of 32 x 32-bit registers
 - Because its registers are 32 bits wide, MIPS is said to have a *word size* of 32 bits
 - This is also the size of the basic unit of transfer between the registers and main memory (although main memory can also be accessed at the granularity of half-words and individual bytes)
 - Processors also have **special registers** such as the *program counter* (PC) and the *stack pointer* (SP)...

How your code is executed



What is your favourite OS?



What do you expect your OS to do for you?

- A program that acts as an intermediary between a user of a computer and the computer hardware (through the ISA)
- Goals:
 - Make the computer system usable
 - User to be able to execute programs as per needs
 - Use of computer hardware efficiently

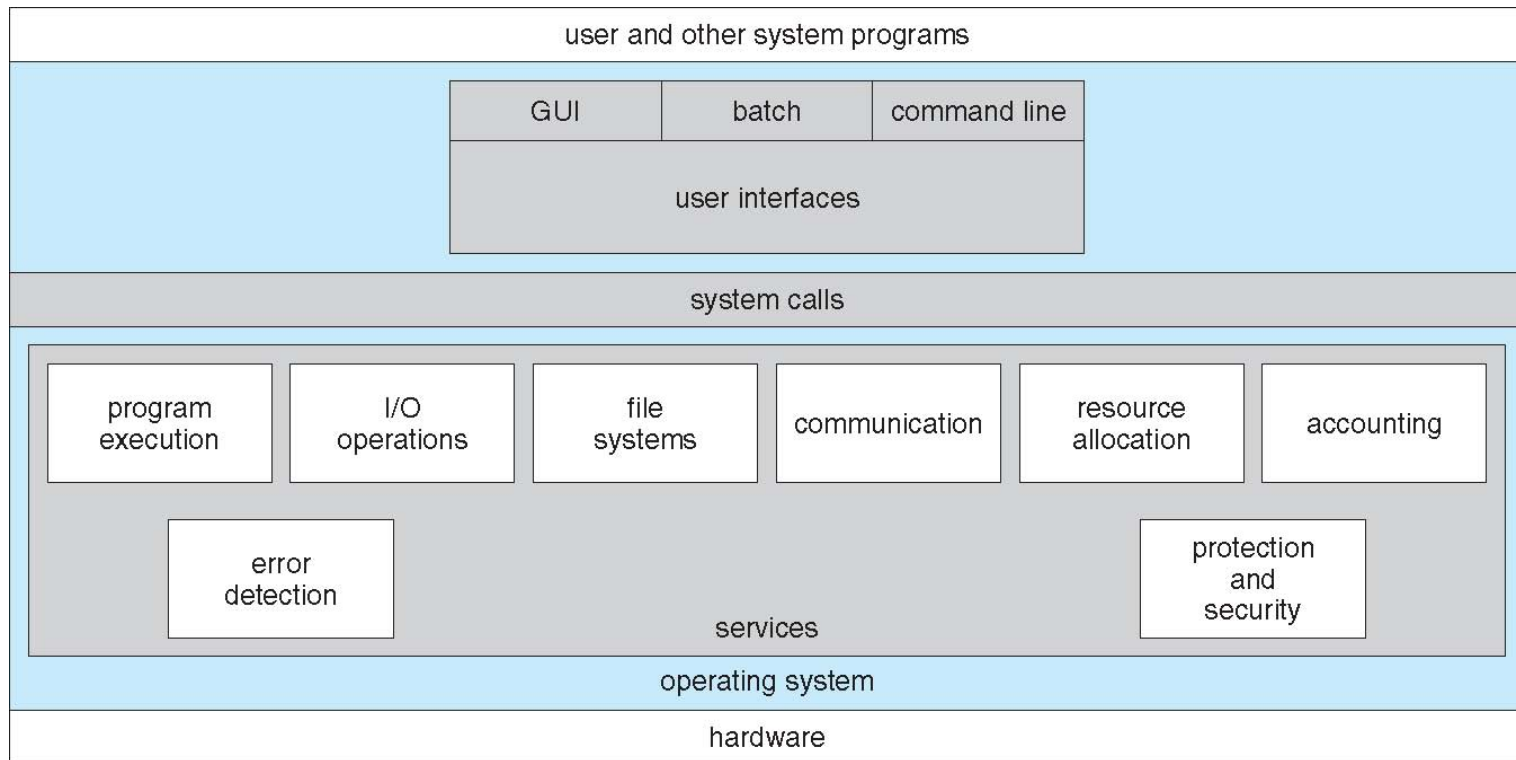
Operating System Roles

- OS acts as a **resource allocator**:
 - Manages all resources
 - In case of conflicting requests ensures efficient and fair resource use
- OS acts as a **control program**:
 - Try to prevent errors and improper use of the computer system

Core components

- Kernel
 - “The program running at all times on the computer; core part of the Operating System”
 - Everything else is either a system program or an application program
- Bootstrap program
 - Also known as firmware, BIOS, etc.
 - Loaded at power-up or reboot
 - Typically stored in ROM (Read Only Memory)
 - Initializes all aspects of system
 - Loads operating system and starts execution




Operating System Services



Memory/Storage Systems characteristics

- Storage systems organized in hierarchy
 - Speed – Cost – Availability & Ruggedness
- Caching
 - Copying information into faster storage system

Memory/Storage-Device Hierarchy



Aside: Levels of Storage

| Level | 1 | 2 | 3 | 4 | 5 |
|---------------------------|--|-------------------------------|------------------|------------------|------------------|
| Name | registers | cache | main memory | solid state disk | magnetic disk |
| Typical size | < 1 KB | < 16MB | < 64GB | < 1 TB | < 10 TB |
| Implementation technology | custom memory with multiple ports CMOS | on-chip or off-chip CMOS SRAM | CMOS SRAM | flash memory | magnetic disk |
| Access time (ns) | 0.25 - 0.5 | 0.5 - 25 | 80 - 250 | 25,000 - 50,000 | 5,000,000 |
| Bandwidth (MB/sec) | 20,000 - 100,000 | 5,000 - 10,000 | 1,000 - 5,000 | 500 | 20 - 150 |
| Managed by | compiler | hardware | operating system | operating system | operating system |
| Backed by | cache | main memory | disk | disk | disk or tape |

Aside: System Latencies

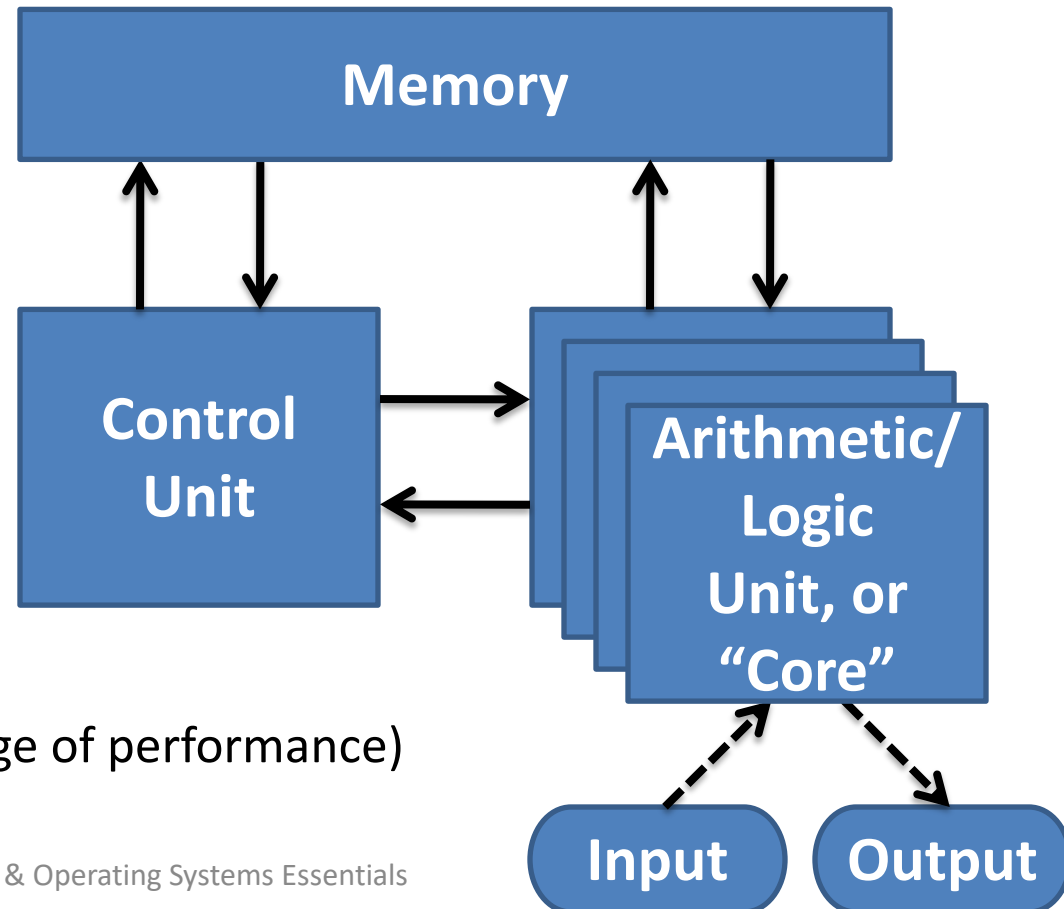
| Event | Latency | Scaled |
|--|-----------|---------------|
| 1 CPU cycle | 0.3 ns | 1 s |
| Level 1 cache access | 0.9 ns | 3 s |
| Level 2 cache access | 2.8 ns | 9 s |
| Level 3 cache access | 12.9 ns | 43 s |
| Main memory access (DRAM, from CPU) | 120 ns | 6 min |
| Solid-state disk I/O (flash memory) | 50–150 µs | 2–6 days |
| Rotational disk I/O | 1–10 ms | 1–12 months |
| Internet: San Francisco to New York | 40 ms | 4 years |
| Internet: San Francisco to United Kingdom | 81 ms | 8 years |
| Internet: San Francisco to Australia | 183 ms | 19 years |
| TCP packet retransmit | 1–3 s | 105–317 years |
| OS virtualization system reboot | 4 s | 423 years |
| SCSI command time-out | 30 s | 3 millennia |
| Hardware (HW) virtualization system reboot | 40 s | 4 millennia |
| Physical system reboot | 5 m | 32 millennia |

As performance demands increase...

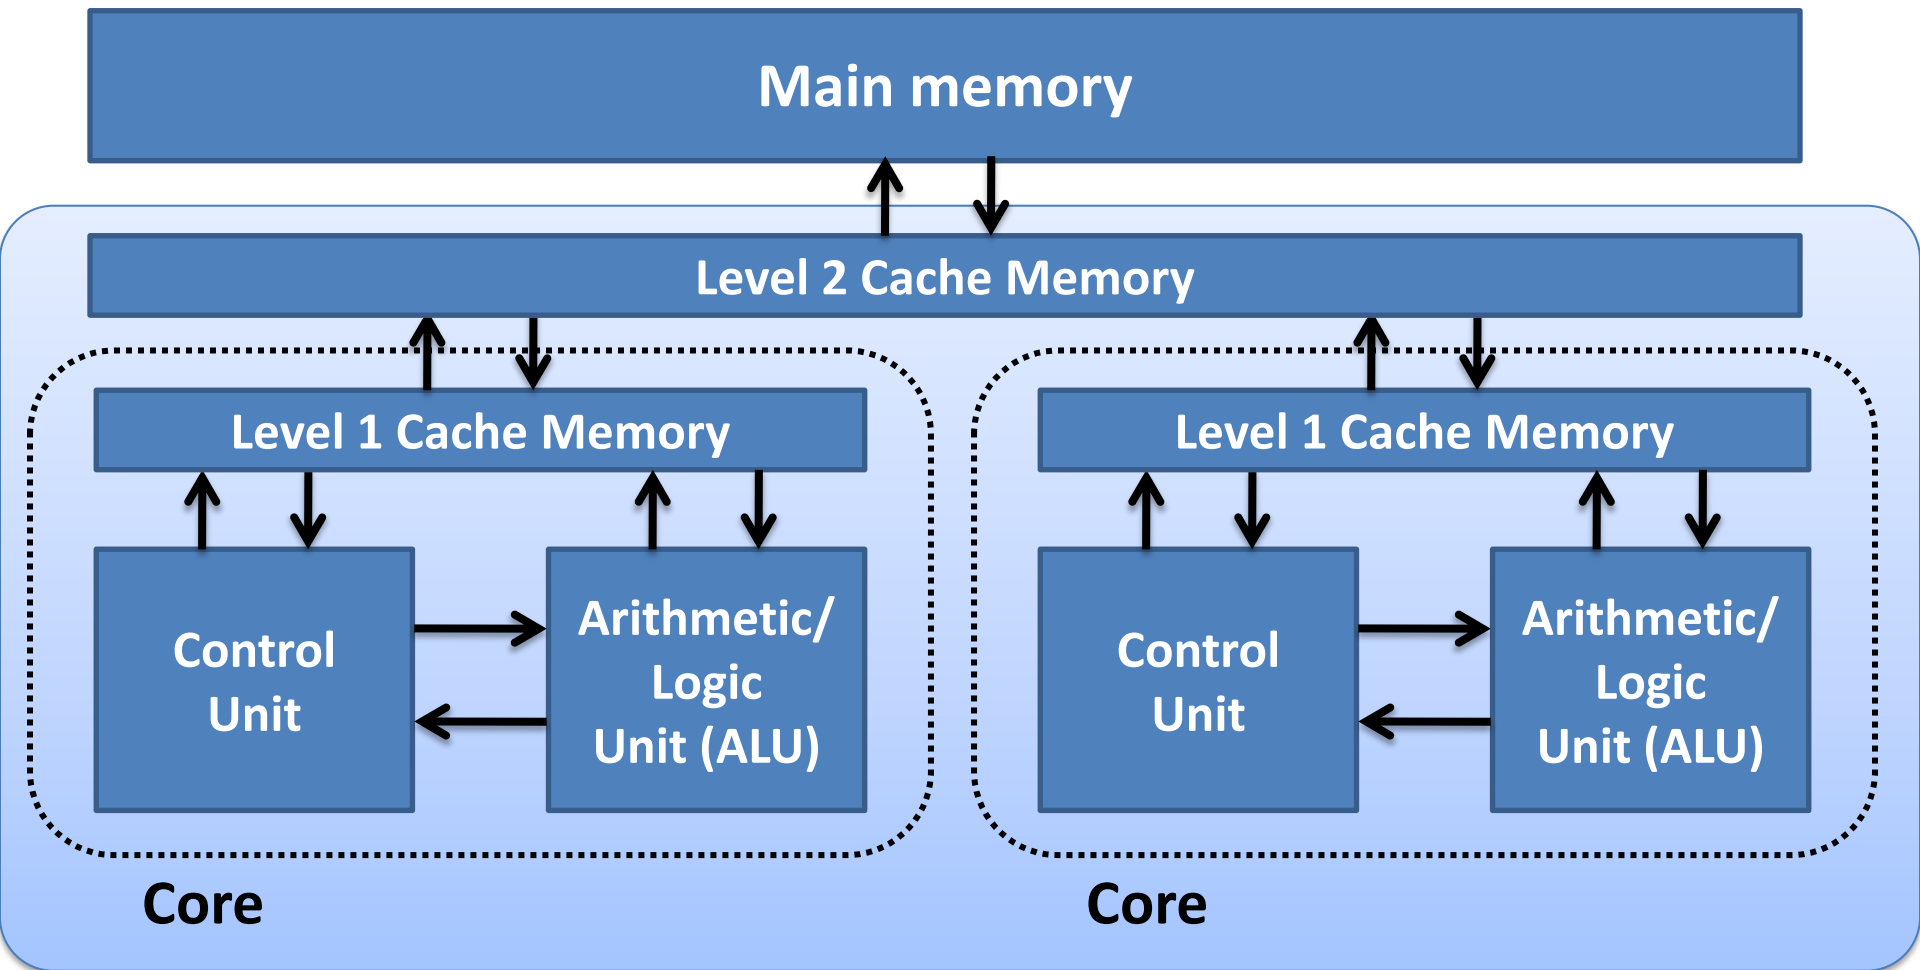
- We see recent renewed interest in *parallel* architectures

– Faster, although

- They complicate system software
- We are not always able to hide the complexity from application software (esp. to take advantage of performance)




Another multi-core architecture



A dual-core processor

Coarser-grained parallelism: clustering

- We can increase performance by linking computers using high-speed networks




- Leads to idea of “blade servers”
 - Obviously they don’t all need screens, etc.
- Applications run across the cluster (ideally)
 - Although, as mentioned earlier, some applications can’t easily be decomposed in this way

Computer System Organisation

- One or more CPUs, device controllers → access to shared memory
- CPUs and devices competing for memory cycles
- Each device controller is in charge of a particular device type and has a local buffer
 - CPU moves data from/to main memory to/from local buffers
 - I/O communication from the device to local buffer of controller.
 - Device controller informs CPU that it has finished its operation by causing an interrupt


Direct Memory Access

- What happens then when an IO request needs to be processed?
 - Wait for it?
- Do something useful: DMA
 - Used by high-speed I/O devices
 - Direct transfer from device controller buffer's storage to main memory region
 - No CPU intervention
 - Use of a single purpose processor
 - DMAC (DMA Controller)



DMA hardware requirements

- With DMA, we have *two* processors accessing main memory, potentially simultaneously
- This requires either **dual-port memory** or **arbitration circuitry**, both of which are quite complex/ expensive



Operating System terms

Multiprogramming, Multitasking, Multiprocessing & Multithreading

Process Management

- A process is a program in execution that needs resources
 - Program is a passive entity, process is an active entity
- Single-threaded process has one program counter specifying location of next instruction to execute
- Multi-threaded process has one program counter per thread

Process Management

- Creating and deleting both user and system processes
- Suspending and resuming processes
- Providing mechanisms for process synchronization
- Providing mechanisms for process communication
- Providing mechanisms for deadlock handling

Memory Management

- All instructions in memory
- Memory management determines what is in memory
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed

Storage Management

- File-System management:
 - Files usually organized into directories
 - Access control on most systems to determine who can access what
- OS:
 - Creating/editing and deleting files and directories
 - Mapping files onto secondary storage
 - Backup files onto stable (non-volatile) storage media

Recommended Reading

- Section 1.3, 1.8 : <https://ict.iitk.ac.in/wp-content/uploads/CS422-Computer-Architecture-ComputerOrganizationAndDesign5thEdition2014.pdf>
- Silberschatz et. al., chapter 1 up to section 1.8