

Week 1: Recap, Binary Radix Sort, Tries

Data Structures:

-

Sorting algorithms:

- Naïve – $O(n^2)$ in the worst/average case
 - Selection sort, insertion sort, bubble sort
- Clever – $O(n \log n)$ in the worst/average case
 - Merge sort, Heapsort
- Fastest in practice – $O(n \log n)$ on average, no better than $O(n^2)$ in the worst case unless a clever variant is used
 - Quicksort

Comparison-based sorting

- **No sorting algorithm based on pairwise comparison of values can be better than $O(n \log n)$**
- Proof in Vid 2

Radix sort

- $O(n)$
- Has to exploit the structure of the items being sorted, so may be less versatile
- In practice, faster than $O(n \log n)$ only for very large n
- Algorithm:
 1. Assume items to sort can be treated as bit-sequences of length m
 2. Let b be a chosen factor of m
 - a. b and m are constants for any particular instance
 3. (Each item has bit positions labelled $0, 1, m-1$ (right to left))
 4. m/b iterations
 - a. In each iteration, the items are distributed into 2^b buckets (lists)
 - i. The buckets are labelled $0, 1, \dots, 2^b - 1$
 - b. During the i^{th} iteration, an item is placed in the bucket corresponding to the integer represented by the bits in positions $b \times i - 1, \dots, b \times i$
 - c. At the end of an iteration, the buckets are concatenated to give a new sequence for the next iteration
- Pseudocode:

```

// assume we have the following method which returns the value
// represented by the b bits of x when starting at position pos
private int bits(Item x, int b, int pos)

// suppose that:
//   a is the sequence to be sorted
//   m is the number of bits in each item of the sequence a
//   b is the 'block length' of radix sort

int numIterations = m/b; // number of iterations required for sorting
int numBuckets = (int) Math.pow(2, b); // number of buckets

// represent sequence a to be sorted as an ArrayList of Items
ArrayList<Item> a = new ArrayList<Item>();

// represent the buckets as an array of ArrayLists
ArrayList<Item>[] buckets = new ArrayList[numBuckets];
for (int i=0; i<numBuckets; i++) buckets[i] = new ArrayList<Item>();

```

○

```

for (int i=1; i<=numIterations; i++){

    // clear the buckets
    for (int j=0; j<numBuckets; j++) buckets[j].clear();

    // distribute the items (in order from the sequence a)
    for (Item x : a){
        // find the value of the b bits starting from position (i-1)*b in x
        int k = bits(x, b, (i-1)*b); // find the correct bucket for item x
        buckets[k].add(x); // add item to this bucket
    }

    a.clear(); // clear the sequence

    // concatenate the buckets (in sequence) to form the new sequence
    for (j=0; j<numBuckets; j++) a.addAll(buckets[j]);
}

```

○

- Correctness
 - [Proof in Vid 3]
- Complexity
 - $O(n)$
 - Number of buckets - 2^b
 - Number of iterations - m/b
 - During each iteration:
 - Bucket allocation - $O(n)$
 - Bucket concatenation - $O(2^b)$
 - Overall: $O\left(\frac{m}{b} \cdot (n + 2^b)\right)$
 - **Time-space trade-off**
 - The larger the value of b , the smaller the multiplicative constant (m/b) in the complexity function and so the faster the algorithm will become
 - However, an array of size 2^b is required for the buckets; therefore, increasing b will increase the space requirements

Tries (retrieval)

- Tries are to binary trees (comparison-based) as Radix sort is to comparison-based sorting
- Basics:
 - Stored items have a key value that is interpreted as a sequence of bits/character
 - Multiway branch at each node, where each branch has an associated symbol, and no 2 siblings have the same symbol
 - The branch taken at level i during a search is determined by the i^{th} element of the key value (i^{th} bit, i^{th} char)

- Tracing a path from the root to a node spells out the key value of the item
- Search pseudocode:

```
// searching for a word w in a trie t
Node n = root of t; // current node (start at root)
int i = 0; // current position in word w (start at beginning)

while (true) {
    if (n has a child c labelled w.charAt(i)) {
        // can match the character of word in the current position
        if (i == w.length()-1) { // end of word
            if (c is an 'intermediate' node) return "absent";
            else return "present";
        }
        else { // not at end of word
            n = c; // move to child node
            i++; // move to next character of word
        }
    }
    else return "absent"; // cannot match current character
}
```

-
- Insert pseudocode:

```
// inserting a word w in a trie t
Node n = root of t; // current node (start at root)

for (int i=0; i < w.length(); i++){ // go through chars of word
    if (n has no child c labelled w.charAt(i)){
        // need to add new node
        create such a child c;
        mark c as intermediate;
    }
    n = c; // move to child node
}
mark n as representing a word;
```

-
- Complexity of operations
 - Almost independent of no. of items
 - Essentially linear in the string length
- Implementation
 - Array (of pointers to represent the children of each node)
 - Linked lists (to represent the children of each node)
 - Time/space trade-off

Week 2: Strings and Text Algorithms

Text compression

- A special case of **data compression**
 - Saves disk space and transmission time
- Must be **lossless** (original must be recoverable without error)
- Some other forms of compression can be lossy (pictures, sound, etc)
- Examples:
 - compress, gzip in Unix, ZIP utilities for Windows
 - 2 main approaches: statistical and dictionary
- **Compression ratio:** x/y
 - x – size of compressed file, y – size of original file
 - measured in B, kB, MB, ...
 - Compressing 10MB to 2 would yield $2/10=0.2$
- **Percentage space saved:** $(1 - \text{"compression ratio"}) * 100\%$
 - Space saved expressed as a percentage of the original file size

- 10 to 2MB -> 80%
- Space savings in the range 40-60% are typical
 - The higher the saving, the better
- Huffman encoding
 - Classical **statistical** method
 - Now mostly superseded in practice by more effective dictionary methods
 - Fixed (ASCII) code replaced by **variable** length code for each char
 - Every char is represented by a unique codeword (bit string)
 - Frequently occurring chars are represented by shorter codewords
 - The code has the **prefix** property
 - No codeword is a prefix of another (gives **unambiguous** decompression)
 - Based on a **Huffman tree** (a proper binary tree)
 - Each char is represented by a leaf node
 - Codeword for a char is given by the path from the root to the appropriate leaf (left=0, right=1)
 - Prefix property follows from this
 - Construction:
 1. Add leaves of Huffman tree
 - a. Characters with their frequencies label the leaf nodes
 2. While there is >1 parentless node:
 - a. Add new parent to nodes of smallest weight
 - b. Weight of new node equals sum of the weights of the child nodes

```
// set up the leaf nodes
for (each distinct character c occurring in the text){
  make a new parentless node n;
  int f = frequency count for c;
  n.setWeight(f); // weight equals the frequency
  n.setCharacter(c); // set character value
  // leaf so no children
  n.setLeftChild(null);
  n.setRightChild(null);
}
// construct the branch nodes and links
while (no. of parentless nodes > 1){
  make a new parentless node z; // new node
  x, y = 2 parentless nodes of minimum weight; // its children
  z.setLeftChild(x); // set x to be the left child of new node
  z.setRightChild(y); // set y to be the right child of new node
  int w = x.getWeight()+y.getWeight(); // calculate weight of node
  z.setWeight(w); // set the weight of the new node
}
```

- Optimality
 - Weighted path length (WPL) of a tree T
 - Sum of weight * distance from root, where sum is over all leaf nodes
 - Huffman tree has minimum WPL over all binary trees with the given leaf weights
 - Huffman tree need not be unique (e.g., nodes > 2 with min weight)
 - However, all Huffman trees for a given set of frequencies have same WPL
 - WPL is the number of bits in compressed file
 - bits = sum over chars (frequency of chars * code length of char)
 - A Huffman tree **minimises** this number

- Hence, Huffman coding is **optimal** for all possible codes built this way
- Algorithmic requirements (n – text length; m – distinct chars in text)
 - Building the tree – **$O(n + m \log m)$**
 - Based on using a heap for the weights
 - $O(n)$ for finding frequencies and then $O(m \log m)$ to build the tree
 - Since m is essentially a constant, it is really **$O(n)$**
 - Compression uses a code table (an array of codes, indexed by char)
 - $O(m \log m)$ to build the table
 - m chars, so m paths of length $O(\log m)$
 - The Huffman tree is also a binary tree, so height is $O(\log m)$
 - $O(n)$ to compress: n chars in the text, so n lookups in array – $O(1)$
 - Overall: **$O(m \log m) + O(n)$**
 - Decompression uses the tree directly (repeatedly trace paths in tree)
 - **$O(n \log m)$** since n is size of text, hence, n paths of length $O(\log m)$
 - Problem – some representation of the Huffman tree must be stored with the compressed file (otherwise, decompression would be impossible)
 - Alternatives:
 - Fixed set of frequencies based on typical values for text
 - Reduces compression ratio
 - Adaptive Huffman coding: the same tree is built and adapted by compressor and by the decompressor as chars are encoded/decoded
 - Slows down compression and decompression (but not by much if clever)
- LZW compression (Lempel, Ziv, and Welch)
 - Popular dictionary-based method
 - The dictionary is a collection of strings
 - Each with a **codeword** that represents it
 - The codeword is a bit string
 - It can be interpreted as a non-negative integer
 - Whenever a codeword is outputted during compression, the **bit string** is written of the compressed file
 - Using a number of bits determined by the **current codeword length**
 - At any point all bit strings are the same length
 - The dictionary is built **dynamically** during compression

```

set current text position i to 0;
initialise codeword length k (say to 8);
initialise the dictionary d;

while (the text t is not exhausted) {

    identify the longest string s, starting at position i of text t
    that is represented in the dictionary d;
    // there is such string, as all strings of length 1 are in d

    output codeword for the string s; // using k bits

    // move to the next position in the text
    i += s.length(); // move forward by the length of string just encoded
    c = character at position i in t; // character in next position

    add string s+c to dictionary d, paired with next available codeword;
    // may have to increment the codeword length k to make this possible
}

```

- (and also during decompression)

```

initialise codeword length k;
initialise the dictionary;

read the first codeword x from the compressed file f; // i.e. read k bits
String s = d.lookup(x); // look up codeword in dictionary
output s; // output decompressed string

while (f is not exhausted){

    String oldS = s.clone(); // copy last string decompressed

    if (d is full) k++; // dictionary full so increase the code word length

    get next codeword x from f; // i.e. read k bits
    s = d.lookup(x); // look up codeword in dictionary
    output s; // output decompressed string

    String newS = oldS + s.charAt(0); // string to add to dictionary
}

```

-
- Initially, dictionary contains all possible strings of length 1
- Throughout execution, the dictionary is **closed under prefixes**
 - i.e., if the string s is represented in the dictionary, so is every prefix of s
- => a **trie** is an ideal representation of the dictionary
 - **every** node in the trie represents a 'word' in the dictionary
 - a trie is effective and efficient ??? [in live lec]
- During (de)compression there is a current **codeword length k**. This value is dynamic
 - Codewords are bit strings of length k
 - Clear how many bits to read when decompressing
 - For a given k, there are exactly 2^k distinct codewords available
 - i.e., all possible bit strings of length k
 - This limits the size of the dictionary
 - However, the codeword length can be incremented when necessary
 - thus doubling the number of available codewords
 - and **all** codewords change (bit strings now of length k+1)
 - Initial value of k should be large enough to encode all strings of length 1
- Variants:
 - Constant codeword length
 - Fix the codeword length for all time
 - The dictionary has fixed capacity: when full, just stop adding to it
 - **Dynamic codeword length**
 - Start with shortest reasonable codeword length, say, 8 for normal text
 - When dictionary becomes full,
 - add 1 to current codeword length (2x no. of codewords)
 - does not affect the sequence of codewords already output
 - May specify a maximum codeword length, as increasing the size indefinitely may become counter-productive
 - LRU (Least Recently Used)
 - When dictionary becomes full and codeword length maximal
 - Current string replaces LRU string in dictionary
- Decompression – Special Case
 - Possible to encounter a codeword that is not (yet) in the dictionary

- because decompression is 'out of phase' with compression
- Possible to deduce what string it must represent
 - Solution: if (lookup fails) $s = \text{oldS} + \text{oldS.charAt}(0)$;
- Complexity
 - Complexity of compression and decompression both **$O(n)$**
 - For a text of length n (if suitably implemented)
 - Algorithms essentially involve just 1 pass through the text
 - Need to "search" for longest strings in the dictionary and then add new strings to the dictionary
 - Efficiently performing this search and adding a string is part of the AE

String comparison

- Notation:
 - $S = s_0 s_1 \dots s_{m-1}$
 - m – length of string
 - $s[i]$ is the $(i+1)$ th element of the string, i.e., s_i
 - $s[i..j]$ – substring from i to j
 - Prefixes and suffixes
 - j -th prefix – the first j characters of s denoted $s[0..j-1]$
 - 0th prefix – empty string
 - j -th suffix – the last j characters of s denoted $s[m-j..m-1]$
 - 0th suffix – empty string
- Fundamental question: how similar (different) are 2 strings?
 - Applications include:
 - biology (DNA and protein sequences)
 - file comparison (diff in Unix)
 - spelling correction, speech recognition, ...
- **Basic operations** for transforming
 - **insert** a char
 - **delete** a char
 - **substitute** a char by another
- String distance
 - **Distance** between s and t – the smallest number of basic operations needed to transform s to t
 - String comparison algorithms use **dynamic programming**
 - Problem solved by building up solutions to sub-problems of ever-increasing size
 - Often called the **tabular method** (builds up a **table** of relevant values)
 - Eventually, one of the values in the table gives the required answer
 - **Dynamic programming**
 - i -th prefix of string s is the first i characters of s
 - Let **$d(i,j)$** be the distance between i -th prefix of s and j -th prefix of t
 - Distance between s and t is then $d(m,n)$ (since lengths of s and t are m and n)
 - Recurrence relation:

$$d(i, j) = \begin{cases} d(i-1, j-1) & \text{if } s[i-1] = t[j-1] \\ 1 + \min\{d(i, j-1), d(i-1, j), d(i-1, j-1)\} & \text{otherwise} \end{cases}$$

no operations are required
the distance is given by that between the $i-1^{\text{th}}$
and $j-1^{\text{th}}$ prefixes of s and t

- - Second line:
 - '1' will perform 1 operation (insertion/deletion/substitution)
 - $d(i, j-1)$ – insertion
 - $d(i-1, j)$ – deletion
 - $d(i-1, j-1)$ – substitution
 - Algorithm comes immediately from the formula
 - Time and space complexity both **$O(mn)$**
 - A consequence of the size of the table
 - Space complexity can easily be reduced to **$O(m+n)$** (by keeping the most recent entry in each column of the table)
 - **Traceback phase** used to construct an **optimal alignment**
 - Trace a path in the table from bottom right to top left
 - Draw an arrow from an entry to the entry that led to its value
 - Can be done using only $O(m+n)$ space with Hirschberg's algorithm
 - Interpretation
 - Vertical step – deletion
 - Horizontal step – insertion
 - Diagonal step – match/substitution
 - A match if the distance doesn't change, substitution otherwise

Week 3: String/Pattern Search

Def:

- Searching a (long) text for a (short) string/pattern
 - Many applications, including:
 - info retrieval
 - text editing
 - computational biology
 - Many variants, such as exact/approximate matches
 - 1st occurrence or all occurrences
 - 1 text and many string/patterns
 - many texts and 1 string/pattern
 - Abstract view
 - given a text t (of length n) and a string/pattern s (of length m), find the position of the first occurrence (if it exists) of s in t
 - usually n is large and m is small

- Algorithms:

- Brute force (exhaustive search)

- Tests all possible positions
 - Wordy:
 - Set the current starting position in text as 0
 - Compare text and string chars left-to-right until the entire string is matched/character mismatches
 - If mismatch:
 - Advance the starting position by 1 and repeat
 - Continue until a match or text is exhausted
 - Uses char arrays in Java (not strings)

```

/** return smallest k such that s occurs in t starting at position k */
public int bruteForce (char[] s, char[] t){
    int m = s.length; // length of string/pattern
    int n = t.length; // length of text
    int sp = 0; // starting position in text t
    int i = 0; // curr position in text
    int j = 0; // curr position in string/pattern s
    while (sp <= n-m && j < m) { // not reached end of text/string
        if (t[i] == s[j]){ // chars match
            i++; // move on in text
            j++; // move on in string/pattern
        } else { // a mismatch
            j = 0; // start again in string
            sp++; // advance starting position
            i = sp; // back up in text to new starting position
        }
    }
    if (j == m) return sp; // occurrence found (reached end of string)
    else return -1; // no occurrence (reached end of text)
}

```

- Complexity
 - Worst case is no better than $O(mn)$
 - Typically, the number of comparisons from each point will be small
 - Often just 1 comparison for mismatch
 - $\Rightarrow O(n)$ on average

- KMP

- Knuth-Morris-Pratt
 - $O(m+n)$ in the worst case
 - **on-line** – removes the need to back-up in the text
 - Involves pre-processing the string to build a border table
 - **Border table**: an array b with entry $b[j]$ for each position j of the string
 - If mismatch at position j in string/pattern:
 - remain on the current text char (do not back-up)
 - The border table shows which string char should next be compares with the current text char
 - Border – a substring that is both a prefix and a suffix (but not the whole string)
 - Some strings don't have borders, then the empty string (length 0) is the longest
 - String search – Border table
 - Border table b of the string pattern

- $b[j]$ = length of the longest border of $s[0...j-1]$

```

/** return smallest k such that s occurs from position k in t or -1 if no k exists */
public int kmp(char[] t, char[] s) {
    int m = s.length; // length of string/pattern
    int n = t.length; // length of text
    int i = 0; // current position in text
    int j = 0; // current position in string s
    int[] b = new int[m]; // create border table
    setUp(b); // set up the border table
    while (i <= n) { // not reached end of text
        if (t[i] == s[j]){ // if positions match
            i++; // move on in text
            j++; // move on in string
            if (j == m) return i - j; // reached end of string so a match
        } else { // mismatch adjust current position in string using the border table
            if (b[j] > 0) // there is a common prefix/suffix
                j = b[j]; // change position in string (position in text unchanged)
            else { // no common prefix/suffix
                if (j == 0) i++; // move forward one position in text if not advanced
                else j = 0; // else start from beginning of the string
            }
        }
    }
    return -1; // no occurrence
}

```

-
- Search – $O(n)$ in worst case
- Border table creation
 - Naïve – $O(m^3)$
 - More efficient: use KMP algorithm $\Rightarrow O(m)$
- Boyer-Moore
 - Almost always **faster** than brute force/KMP
 - Many apps
 - Properties:
 - Typically, many text characters are skipped without even being checked
 - String/pattern is scanned **right-to-left**
 - Text characters involved in a mismatch are **used to decide** next comparison
 - Involves pre-processing the string to record the position of the **last** occurrence of each character c in the alphabet
 - The alphabet must be fixed in advance of the search
 - Last occurrence position of character c in the string s (-1 if not present)
 - Simplified version – Boyer-Moore-Horspool
 - Complexity
 - Worst case – $O(mn)$
 - Using good suffix rule, linear ($O(m+n)$)

Week 4: Graphs and Graph Algorithms

Recap of AF2:

- Undirected graph $G = (V, E)$
 - V – finite set of vertices (vertex set)
 - E – set of edges, each edge is a subset of V of size 2 (edge set)
 - Pictorially

- V – point
 - Edge – line connecting points
 - Representations of same graph can be different
- Properties
 - Vertices are **adjacent** if there is an edge connecting them
 - **Non-adjacent** – not
 - Vertex is **incident** to edge if the edge connects it
 - Path (length = number of edges)
 - Cycle – path that ends where it starts
 - **Degree** – number of edges a vertex is incident to
 - Connected graph if every pair of vertices is joined by a path
 - Non-connected – 2+ connected components
 - **Tree** – connected and acyclic (no cycles) graph
 - Tree with n vertices has n-1 edge
 - **Forest** – acyclic and components are tree
 - **Complete** (a **clique**) – every pair of vertices is joined by an edge
 - **Bipartite** – vertices are in 2 disjoint sets U & W and every edge joins a vertex in U to a vertex in W
- Directed graph (digraph) $D = (V, E)$
 - each edge is an ordered pair (x, y) of vertices
 - “adjacent to” and “adjacent from”
 - Pictorially
 - Edges are drawn as directed lines/arrows
 - Paths and cycles must follow edge directions

Graph representations

- Undirected graphs:
 - Adjacency matrix
 - 1 row and column for each vertex
 - row i, column j contains a 1 if i-th and j-th vertices are adjacent, 0 otherwise
 - Adjacency lists
 - 1 list for each vertex
 - list i contains an entry for j if the vertices i and j are adjacent
- Directed graphs
 - Adjacency matrix
 - 1 row, col for each vertex
 - row i, col j contains 1 if there is an edge from i to j, 0 otherwise
 - Adjacency lists
 - 1 list for each vertex
 - list for vertex i contains vertex j if there is an edge from i to j

Implementation – Adjacency lists

- Define classes for:
 - entries of adjacency lists

```

/** class to represent an entry in the adjacency list of a vertex
in a graph */
public class AdjListNode {

    private int vertexIndex; // the vertex index of the entry

    // possibly other fields, for example representing properties
    // of the edge such as weight, capacity, ...

    /** creates a new entry for vertex indexed i */
    public AdjListNode(int i){
        vertexIndex = i;
    }
    public int getVertexIndex(){ // gets the vertex index of the entry
        return vertexIndex;
    }
    public void setVertexIndex(int i){ // sets vertex index to i
        vertexIndex = i;
    }
}

```

- vertices (includes linked list representing its adjacency list)

```

import java.util.LinkedList; // we require the linked list class
/** class to represent a vertex in a graph */
public class Vertex {

    private int index; // the index of this vertex
    private LinkedList<AdjListNode> adjList; // the adjacency list of vertex

    // possibly other fields, e.g. representing data stored at the node

    /** create a new instance of vertex with index i */
    public Vertex(int i) {
        index = i; // set index
        adjList = new LinkedList<AdjListNode>(); // create empty adjacency list
    }

    /** return the index of the vertex */
    public int getIndex(){
        return index;
    }
}

```

```

// class Vertex continued

/** set the index of the vertex */
public void setIndex(int i){
    index = i;
}

/** return the adjacency list of the vertex */
public LinkedList<AdjListNode> getAdjList(){
    return adjList;
}

/** add vertex with index j to the adjacency list */
public void addToAdjList(int j){
    adjList.addLast(new AdjListNode(j));
}

/** return the degree of the vertex */
public int vertexDegree(){
    return adjList.size();
}
}

```

- graphs (includes size of graph and an array of vertices)
 - array allows for efficient access of “index” of vertex

```

import java.util.LinkedList; // again require the linked list class
// (to add graph algorithms we will need to access adjacency lists)
/** class to represent a graph */
public class Graph {

    private Vertex[] vertices; // the vertices
    private int numVertices = 0; // number of vertices
    // possibly other fields representing properties of the graph

    /** Create a Graph with n vertices indexed 0,...,n-1 */
    public Graph(int n) {
        numVertices = n;
        vertices = new Vertex[n];
        for (int i = 0; i < n; i++) vertices[i] = new Vertex(i);
    }

    /** returns number of vertices in the graph */
    public int size(){
        return numVertices;
    }
}

```

Graph search and traversal algorithms

- A systematic way to explore a graph (when starting from some vertex)
- Applications: web crawler with HTML
- Search/traversal visits all vertices by travelling along edges
 - Traversal is **efficient** if it explores graph in $O(|V|+|E|)$ time
- **Depth-First Search (DFS)**
 - From starting index,
 - Follow a path of unvisited vertices **until path can be extended no further**
 - Then backtrack along the path until an unvisited vertex can be reached
 - Continue until cannot find any unvisited vertices
 - Edges traversed form a spanning tree/forest

- Depth-first spanning tree (forest)
- Spanning tree of a graph is a tree composed of all vertices and some (or all) of the edges of the graph

- Implementation

- Add to vertex class

```
private boolean visited; // has vertex been visited in a traversal?
private int pred; // index of the predecessor vertex in a traversal

public boolean getVisited(){
    return visited;
}
public void setVisited(boolean b){
    visited = b;
}
public int getPred(){
    return pred;
}
public void setPred(int i){
    pred = i;
}
```

- Add to graph class

```
/** visit vertex v, with predecessor index p, during a dfs */
private void visit(Vertex v, int p){
    v.setVisited(true); // update as now visited
    v.setPred(p); // set predecessor (indicates edge used to find vertex)
    LinkedList<AdjListNode> L = v.getAdjList(); // get adjacency list

    for (AdjListNode node : L) // go through all adjacent vertices
        int i = node.getIndex(); // find index of current vertex in list
        if (!vertices[i].getVisited()) // if vertex has not been visited
            visit(vertices[i], v.getIndex()); // continue dfs search from it
            // setting the predecessor vertex index to the index of v
    }
}

/** carry out a depth first search/traversal of the graph */
public void dfs(){
    for (Vertex v : vertices) v.setVisited(false); // initialise
    for (Vertex v : vertices) if (!v.getVisited()) visit(v, -1);
    // if vertex is not yet visited, then start dfs on vertex
    // -1 is used to indicate v was not found through an edge of the graph
}
```

- Complexity

- Overall $O(n+m)$
 - n – number of vertices
 - m – number of edges
- Can be adapted to the adjacency matrix representation
 - Then it's $O(n^2)$ to look at every entry
- Some apps:
 - Determine if a given graph is connected
 - Identify the connected components of a graph
 - Determine if a given graph contains a cycle
 - Determine if a given graph is bipartite

- **Breadth-First Search (BFS)**

- Search fans out as widely as possible at each vertex
 - From the current vertex, visit all the adjacent vertices (**processing** the current vertex)
 - Vertices are processed in the order in which they are visited (visited vertices added to/removed from a queue)
 - Continue until all vertices in current component have been processed
 - Repeat for other components
- Edges traversed form a spanning tree (forest)
 - Breadth-first spanning tree (forest)
- Implementation
 - [AE]
- Complexity

- Overall $O(n+m)$
 - n – number of vertices
 - m – edges
 - Each vertex is visited and queued exactly once
- Example application: Finding **distance** between 2 vertices (v and w)
 - Assign distance to v to be 0
 - Carry out BFS from v
 - When visiting a new vertex, assign distance++
 - Stop when w is reached

Weighted graphs

- Each edge has weight by $wt(e) > 0$
 - Graph undirected or directed
 - Represents lots of things (distance, cost)
 - If an edge is not part of the graph, its weight is infinity
- Representation
 - Adjacency matrix becomes weight matrix
 - Adjacency lists include weight in node
- Shortest paths
 - Length of a span – sum of weights of edges
- Dijkstra's algorithm
 - in 20 minutes, 26 y/o
 - Shortest path between 1 vertex (u) and all others
 - Maintains a set S containing all vertices for which shortest path with u is currently known
 - S initially contains only u
 - Each vertex v has a label $d(v)$ indicating length of shortest path from u to v passing only through vertices in S
 - If no path exists, $d(v) = \text{infinity}$
 - If v is in S , then $d(v)$ is length of shortest path between u and v
 - Invariant of the algorithm: if v is in S and w is not, then the length of the shortest path between u and w is at least that between u and v
 - At each step we add to S the vertex v not in S such that $d(v)$ is minimum
 - After adding v to S , carry out edge relaxation operations (update length $d(w)$ for all vertices w still not in S)
 - Edge relaxation
 - Suppose v and w are not in S , then we know
 - Shortest path between u and v passing only through S equals $d(v)$
 - shortest path between u and w passing only through S equals $d(w)$
 - Suppose v added to S and edge $e = \{v, w\}$ has weight $wt(e)$
 - Calculate shortest path between u and w passing only through $S \cup \{v\}$
 - Shortest path is either:
 - original path through S of length $d(w)$, or
 - path combining e and shortest path between v and u which has length $wt(e) + d(v)$
 - $d(w) = \min(d(w), d(v) + wt(e))$
 - Implementation

-

- An example of a problem in **combinatorial optimisation**
 - find best way of doing something among a (large) number of candidates

- can always be solved, at least in theory, by exhaustive search
 - may be infeasible in practice
 - typically an exponential-time algorithm
 - e.g., K_n (clique of size n) has n^{n-2} spanning trees (Cayley's formula)
 - graph is a clique if every pair vertices is joined by an edge
- Prim-Jarnik minimum spanning tree algorithm
 - example of a greedy algorithm
 - makes a sequence of decisions based on local optimality
 - ends up with the globally optimal solution
- For many problems, greedy algorithms do not yield optimal solution

Prim-Jarnik algorithm

- Pseudocode:

```

set an arbitrary vertex r to be a tree-vertex (tv);
set all other vertices to be non-tree-vertices (ntv);
while (number of ntv > 0){
  find edge e = {p,q} of graph such that
    p is a tv;
    q is an ntv;
    wt(e) is minimised over such edges;
  adjoin edge e to the (spanning) tree;
  make q a tv;
}

```

-
- Analysis:
 - initialisation $O(n)$
 - outer loop $n-1$
 - inner loop $O(n^2)$ (all edges from a tree-vertex to a non-tree-vertex)
 - updating tree $O(1)$
 - Overall: $O(n^3)$
- Correctness:
 - [proof not part of the exam]
 - Yes
- Dijkstra's Refinement:
 - Introduce an attribute bestTV for each non-tree vertex q
 - bestTV is set to the tv p for which $wt(\{p,q\})$ is minimised
 - Pseudocode:

```

set an arbitrary vertex r to be a tree-vertex (tv);
set all other vertices to be non-tree-vertices (ntv);
for (each ntv s) set s.bestTV = r; // r is the only tv

while (number of ntv > 0){
  find ntv q for which wt({q, q.bestTV}) is minimal;
  adjoin {q, q.bestTV} to the tree;
  make q a tv;

  for (each ntv s) update s.bestTV;
  // update bestTV as tree vertices have changed
}

```

- Analysis:
 - initialisation $O(n)$
 - while loop $n-1$
 - first part in while $O(n)$
 - $O(n)$ to find minimal ntv and $O(1)$ to ???
 - second part $O(n)$
 - Overall: $O(n^2)$

Topological ordering

- **Directed Acyclic Graph (DAG)** – directed graph with no cycles
- **Topological order** on a DAG is a labelling of the vertices $1, \dots, n$ such that $(u, v) \in E$ implies $\text{label}(u) < \text{label}(v)$
 - Many apps: scheduling, PERT networks, deadlock detection
 - Scheduling and PERT networks:
 - can model a project with a DAG
 - vertices are tasks/activities, edges indicate dependencies
 - if topological order, can find longest path or longest weighted path
 - can then determine which activities are “critical” (i.e., on longest paths)
- A directed graph D has a topological order if and only if it is a DAG
- **Source** – vertex of in-degree 0 and a sink has out-degree 0
- **DAG has at least 1 source and at least 1 sink**
 - forms the basis of a topological ordering algorithm
 - if there is no source/sink, can build a cycle => not acyclic
 - if no source or sink, can always keep adding vertices to the start/end of a path respectively
 - eventually must add the same vertex twice to the path as there are only finitely many vertices => create a cycle
- Algorithm
 - Add 2 integer attributes to every vertex in the graph:
 - label
 - label in the order
 - count
 - initially equals the number of incoming edges (in-degree) of the vertex
 - updated as the algorithm labels vertices
 - always equals the number of incoming edges from vertices **not labelled**
 - require the label of this vertex $>$ all incoming vertices
 - if all vertices that have incoming edges have been labelled, can just label this vertex with a greater value
 - when attribute becomes 0, add vertex to a queue to be labelled
 - any source vertex can be added to the queue immediately
 - Pseudocode:

```

// assume each vertex has 2 integer attributes: label and count
// count is the number of incoming edges from unlabelled vertices
// label will give the topological ordering

for (each vertex v) v.setCount(v.getInDegree()); // initial count values

Set up an empty sourceQueue

for (each vertex v) // add vertices with no incoming edges to the queue
  if (v.getCount() == 0) add v to sourceQueue; // i.e. source vertices

int nextLabel = 1; // initialise labelling (gives topological ordering)
while (sourceQueue is non-empty){
  dequeue v from sourceQueue;
  v.setLabel(nextLabel++); // label vertex (and increment nextLabel)
  for (each w adjacent from v){ // consider each vertex w adjacent from v
    w.setCount(w.getCount() - 1); // update attribute count
    // add vertex to source queue if there are no incoming vertices
    if (w.getCount() == 0) add w to sourceQueue;
  }
}

```

- Correctness
 - A vertex is given a label only when the number of incoming edges from unlabelled vertices is 0
 - all predecessor vertices must already be labelled with smaller numbers
 - dependent on using a queue (FIFO for labelling)
- Analysis (n vertices, m edges)
 - for adjacency lists representation
 - finding in-degree of each vertex is $O(n_m)$
 - main loop is n times
 - every list is scanned again once
 - Overall: $O(n+m)$
 - for adjacency matrix representation
 - Overall: $O(n^2)$
- Deadlock detection
 - Method 1: an adaptation of the topological ordering algorithm
 - if the source queue becomes empty before all vertices are labelled, then there must be a cycle
 - if all vertices can be labelled, then the digraph is acyclic
 - Method 2: an adaptation of DFS
 - when a vertex u is 'visited', check whether there is an edge from u to a vertex v which is on the current path from the current starting vertex
 - the existence of such a vertex indicates a cycle (adaptation of DFS since need to 'remember' current path)

Week 6: Intro to NP Completeness

Polynomial vs exponential time:

- Exponential-time algorithms are in general "bad"
 - increases in processor speeds do not lead to significant changes in the slow behaviour when the input size is large
- Polynomial-time algorithms are in general "good"
- "Efficient algorithms" are in polynomial time
 - Require extra insight
 - Exponential are usually exhaustive
- **Polynomial-time solvable** problem – if it admits a polynomial-time algorithm
 - **Intractable** problem – (proved to be) impossible to find such algorithm
 - NP-complete problem – difficult problem which no one has found a solution for yet (but hasn't proved to be impossible)

NP-complete problems

- No polynomial-time algorithm is known for an NP-complete problem
 - one of them is solvable in polynomial time => they all are
- No proof of intractability is known for an NP-complete problem

- one of them is intractable => they all are
- Strong belief that NP-complete problems are intractable

Intractable problems

- Causes of intractability:
 - Polynomial time is not sufficient to discover solution
 - Solution itself is so large that exponential time is needed to output it
- Terms:
 - **Undecidable** – no algorithm could solve it
 - Some decidable problems are intractable (can be solved, but not in polynomial time)
- Examples:
 - Roadblock

NP-complete problems

- Problem – characterised by (unspecified) parameters
 - Typically, infinitely many instances for a given problem
- Problem instance – created by giving values to parameters
- Examples:
 - Hamiltonian Cycle (HC)
 - Instance: graph G
 - Question: does G contain a cycle that visits each vertex exactly once?
 - Decision problem – yes/no
 - Every instance is a yes-instance or no-instance
 - Travelling Salesman Decision Problem (TSDP)
 - Instance: a set of n cities and integer distance $d(i, j)$ between each pair of cities i, j , and a target integer K
 - Question: Is there a 'travelling salesman tour' of length $\leq K$
 - Clique Problem (CP)
 - Instance: a graph G and target int K
 - Question: does G contain a clique of size K ?
 - A set of K vertices for which there is an edge between all pairs
 - Ex:
 - There is a clique of size 4, but none of size 5
 - Graph Colouring Problem (GCP)
 - Instance: a graph G and target int K
 - Question: Can one of K colours be attached to each vertex of G so that adjacent vertices always have different colours?
 - Satisfiability (SAT)
 - Instance: Boolean expression B in conjunctive normal form (CNF)
 - Question: is B satisfiable?
 - Can values be assigned to the variables that make B true?
- Optimisation and search problems
 - Optimisation problem – find max/min value
 - e.g., the travelling salesman optimisation problem is to find the min length of tour
 - Search problem – find some appropriate optimal structure
 - e.g., the travelling salesman search problem is to find a min length tour

- NP-completeness deals primarily with decision problems
 - Corresponding to each instance of an optimisation/search problem is a family of instances of a decision problem by setting 'target' values
 - Almost invariably, an optimisation/search problem can be solved in polynomial time if and only if the corresponding decision problem can

The class P

- P – class of all decision problems that can be solved in polynomial time
- Ex:
 - is there a path of length $\leq K$ from vertex u to vertex v in a graph G ?
 - is there a spanning tree of weight $\leq K$ in a graph G ?
 - is a graph G bipartite?
 - is a graph G connected?
 - deadlock detection: does a directed graph D contain a cycle?
 - text searching: does a text t contain an occurrence of a string s ?
 - string distance: is $d(s, t) \leq K$ for strings s and t ?
- Often extended to include search and optimisation problems

The class NP

- The decision problems solvable in **non-deterministic** polynomial time
 - Non-deterministic algorithm can make non-deterministic choices
 - It's allowed to guess
 - Hence, it's **apparently** more powerful than a normal deterministic algorithm
- P is contained within NP
 - A deterministic algorithm is a special case of a non-deterministic one
 - No problem known to be in NP and known not to be in P
- Relationship between P and NP is most notorious unsolved question (million-dollar prize)

Non-deterministic algorithms (NDAs)

- Such an algorithm has an extra operation: non-deterministic choice
- NDA **"solves"** a decision problem if:
 - for a yes-instance, there is **some** execution that returns yes
 - for a no-instance, there is **no** execution that returns yes
- "solves" a decision problem **in polynomial time** if:
 - for every yes-instance, there is **some** execution that returns 'yes' which uses a number of steps bounded by a polynomial in the input
 - for a no-instance, there is no execution that returns yes
- Not useful in practice
 - But a useful mathematical concept for defining the classes of NP and NP-complete problems
- Ex:
 - Graph Colouring

```

// return true if graph g is k-colourable and false otherwise
boolean nDGC(Graph g, int k){

    for (each vertex v in g) v.setColour(nonDeterministicChoice(k));

    for (each edge {u,v} in g)
        if (u.getColour() == v.getColour()) return false;
    return true;
}


```

"verify" the
colouring

"guess" a colour
for each vertex

- Structure:

- Guessing stage (non-deterministic)
- Checking stage (deterministic and polynomial time)




Polynomial time reductions

- PTR – mapping f from a decision problem to another decision problem such that
 - for every instance of first DP
 - the instance $f(\text{first instance})$ of second DP can be constructed in polynomial time
 - $f(\text{first instance})$ is a yes-instance of second DP if and only if first instance is a yes-instance of first DP
- There is a polynomial-time reduction from Π_1 to Π_2 – (cut infinity sign)
- Properties
 - Transitivity
 - Π_1 to Π_2 and Π_2 to Π_3 implies Π_1 to Π_3
 - [proof in lecture]
 - Relevance to P: Π_1 to Π_2 and Π_2 in P implies that Π_1 in P
 - [proof in lecture]

Formal def of NP-completeness

- A decision problem Π is NP-complete if:
 - $\Pi \in NP$
 - For **every** problem $\Pi' \in NP$: Π' is polynomial-time reducible to Π
- Consequences:
 - If Π is NP-complete and Π in P, then $P=NP$
 - Every problem in NP can be solved in polynomial time by reduction to Π
 - Supposing $P \neq NP$, if Π is NP-complete, then Π is not in P

The structure of NP if $P \neq NP$



How to prove a problem is NP-complete

- Suppose know 1 NP-complete problem Π_1
- To prove Π_2 is NP-complete, enough to show:
 - Π_2 is in NP
 - There exists a PTR from Π_1 to Π_2
- Correctness of approach:
 - For any Π in NP, since Π_1 is NP-complete, we have Π PTR to Π_1
 - Since Π PTR to Π_1 and Π_1 PTR to Π_2 , it follows that Π PTR to Π_2

Cook's Theorem (1971): Satisfiability (SAT) is NP-complete

- the proof consists of a generic PTR to SAT from an abstract definition of a general problem in the class NP
- the generic reduction could be instantiated to give an actual reduction for each individual NP problem
- Given Cook's theorem, to prove a decision problem Π is NP-complete, it is sufficient to show:
 - Π is in NP
 - There exists a PTR from SAT to Π

Problem restrictions

- **Restriction** – consists of a subset of the instance of the original problem
 - If a restriction of a given DP Π is NP-complete, then so is Π
 - Given NP-complete problem Π , a restriction of Π **might** be NP-complete
- Examples:
 - Clique restricted to cubic graphs is in P
 - Graph colouring restricted to cubic graphs is NP-complete
 - K-colouring
 - 2-colouring is in P
 - 3-colouring is NP-complete
 - K-SAT
 - 2-SAT is in P
 - 3-SAT is NP-complete (easy to show that $SAT \propto 3-SAT$)
 - [proof in lecture]

Coping with NP-completeness

- Investigate only a **restricted** version of interest (which maybe is in P)
- Seek an exponential-time algorithm improving on exhaustive search
 - e.g., **backtracking**, **branch-and-bound**

- Should extend the set of solvable instances
- For an optimisation problem,
 - settle for an **approximation algorithm** that runs in polynomial time
 - especially if it gives a provably good result (within some factor of optimal)
 - use a **heuristic**
 - e.g., **genetic algorithms, simulated annealing, neural networks**
- For a decision problem,
 - settle for a **probabilistic** algorithm (correct answer with high probability)

Week 7: Computability, Finite-State Automata

Introduction

- Computer?
 - input $x \rightarrow$ black box \rightarrow output $f(x)$
- Computability concerns which **functions** can be computed
 - “What problems can(not) be solved by a computer?”
 - Need for formal definition
- Unsolvable problems
 - Cannot be solved by a computer even with unbounded time
 - Ex: Tiling Problem
 - Instance: a finite set S of tile descriptions
 - Question: can any finite area, of any size, be completely covered using only tiles of types in S , so that adjacent tiles colour match?
 - No algorithm for this
- Undecidable problems
 - A problem Π that admits no algorithm is called **non-computable (unsolvable)**
 - If Π is a decision problem and Π admits no algorithm, it is called **undecidable**
 - e.g., tiling problem
 - Ex: Post’s correspondence problem (PCP)
 - Instance 2 finite sequences of words
 - Question: Does concatenating the words in a given sequence give the same results?

The Halting Problem

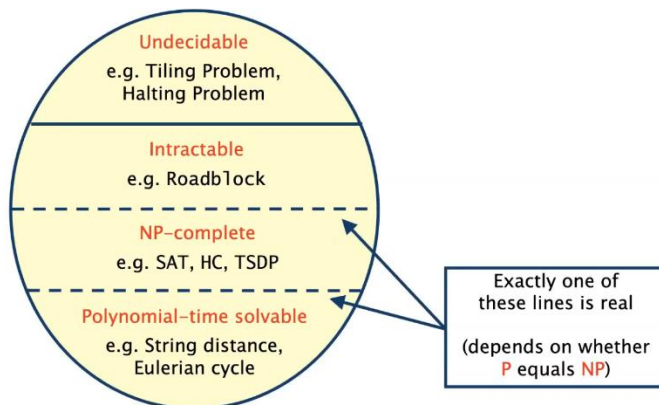
- Impossible project:
 - Write a program Q that takes as input
 - a legal program X
 - an input string S for program X
 - and returns as output
 - yes if program X halts when run with input S
 - no if program X enters an infinite loop when run with input S
- Undecidable
 - Proof
- Formal definition

- Instance: a legal program X and an input string S for X
- Question: does X halt when run with input S ?
- Theorem: HP is undecidable.
- Proof by contradiction:
 - [proof in lecture]

Proving undecidability by reduction:

- [proof in lecture]

Hierarchy of decision problems:



Models of computation

- (Attempts to define the “black box”)
- **Finite-State Automata (FA)**
 - simple machines with a fixed amount of memory
 - have very limited (but still useful) problem-solving ability
- **Pushdown Automata (PDA)**
 - simple machines with an unlimited memory that behaves like a stack
- **Turing Machines (TM)**
 - simple machines with an unlimited memory that can be used essentially arbitrarily
 - essentially the same power as a typical computer

Deterministic finite-state automata (DFA)

- Simple machines with limited memory which **recognise** input on a read-only tape
- Consists of:
 - a finite input alphabet Σ
 - a finite set of states Q
 - an initial/start state $q_0 \in Q$ and a set of accepting states $F \subseteq Q$
 - control/program or transition relation $T \subseteq (Q \times \Sigma) \times Q$
- Defines a **language**
 - determines whether the string on the input tape belongs to that language
 - in other words, it solves a decision problem (yes if belongs, no if doesn't)
- Recognises/accepts a language
 - the input strings which, when 'run', end in an accepting state

Non-deterministic finite-state automata (NFA)

- Recognition is similar to non-deterministic algorithms “solving” a DP
 - Only require there exists a ‘run’ that ends in an accepting state
 - i.e., under 1 possible resolution of the non-deterministic choices
- Any NFA can be converted into a DFA (reduction)
 - \Rightarrow Non-determinism does not expand the class of languages that can be recognised by FA
 - being able to guess does not give it any extra power
 - states of the DFA are sets of states of the NFA
 - construction can cause a blow-up in the number of states
 - in worst case, from N states to 2^N states

Regular languages and regular expressions

- **Regular languages** – languages that can be recognised by FA
- A regular language (over an alphabet Σ) can be specified by a **regular expression** over Σ
 - ϵ (empty string) is an RE
 - σ is an RE
- If R and S are REs, then so are:
 - RS (concatenated)
 - $R \mid S$ (choice between R or S)
 - R^* (0 or more copies of R) (**closure**)
 - (R) (override preference with brackets)

Regular expressions

- Order of precedence (highest first):
 - closure ($*$), concatenation, choice (\mid)
 - brackets to override
- Additional operations:
 - Complement $\neg x$
 - equivalent to the ‘or’ of all characters except x
 - Any single character ?
 - equivalent to the ‘or’ of all characters
- Example:
 - No DFA that can recognise strings of the form $a^n b^n$ (number of a ’s followed by same number of b ’s)
 - Can be done with some form of memory, e.g., a stack

There are some functions (languages) that are computable, but not by an FA

Week 8: Pushdown Automata

Non-deterministic PDA

- Consists of:
 - a finite input alphabet Σ , a finite set of stack symbols G

- a finite set of states Q including start state and set of accepting states
 - control or transition relation $T \subseteq (Q \times \Sigma \cup \{\varepsilon\} \times G \cup \{\varepsilon\}) \times (Q \times G \cup \{\varepsilon\})$
 - ε – empty set
- Accepts an input if and only if, after the input has been read, the stack is empty and control is in an accepting state
- There is no explicit test that the stack is empty
 - This can be achieved by adding a special symbol ($\$$) to the stack at the start of the computation
 - i.e., we add the symbol to the stack when we know the stack is empty and never add $\$$ at any other point
 - Check for emptiness by checking if $\$$ is on top
- Deterministic PDAs (**DPDAs**) are **less powerful**
 - There are languages that can be recognised by NDPDA but not by DPDA (e.g., language of palindromes)
- **Palindromes**
 - How to recognise:
 - Push first half onto stack
 - As we read each new char, check it is the same as top element and pop this
 - Enter an accepting state if all checks succeed
 - Why **non-determinism**?
 - Need to “guess” where the **middle** is
 - and if there’s even/odd chars
 - Cannot work this out first and then check the string, would need:
 - to read string twice
 - an unbounded number of states if string is infinite
 - Automaton recognises $\{ a^n b^n \mid n \geq 0 \}$
 - \Rightarrow PDAs are more powerful than FAs
- Languages that can be recognised by a PDA are **context-free languages**
 - Are all languages context-free?
 - No, consider $a^n b^n c^n$

Turing machines

- T to recognise a particular language consists of:
 - a finite alphabet Σ , including a blank symbol ($\#$)
 - an unbounded tape of squares
 - each can hold a single symbol
 - tape unbounded in both directions
 - a tape head that scans a single square
 - can read from it and write to the square
 - then moves one square left/right along the tape
 - a set S of states
 - includes a single start state s_0 and 2 halt (terminal) states s_Y and s_N
- Transition function:
 - $f: (S \setminus \{s_Y, s_N\}) \times \Sigma \rightarrow (S \times \Sigma \times \{Left, Right\})$
 - For each non-terminal state and symbol, the function f specifies:
 - a new state (perhaps unchanged)
 - a new symbol (perhaps unchanged)

- a direction to move along the tape
 - $f(s, \sigma)(s', \sigma', d)$ means reading symbol σ from the tape in state s
 - move to state $s' \in S$
 - overwrite the symbol σ on the tape with the symbol $\sigma' \in \Sigma$.
 - move the tape head one square in direction d (left/right)
- Computation
 - The (finite) input string is placed on the tape
 - assume initially all other squares of the tape contain blanks
 - The tape head is placed on the first symbol of the input
 - T starts in state s_0 (scanning the first symbol)
 - If T halts in s_Y , 'yes'
 - Otherwise, 'no'
- Palindrome problem
 - Instance: a finite string Y
 - Question: is Y a palindrome, i.e., is Y equal to the reverse of itself?
 - [lecture]
- A TM can be described by its state transition diagram – directed graph where:
 - each state is a vertex
 - ...
- The TM that accepts language L actually computes the function f where $f(x)=1$ if x is in L , 0 otherwise
 - Definition of a TM can be amended as follows:
 - to have a set H of halt states
 - the function it computes is defined by $f(x)=y$ where:
 - x – initial string on tape
 - y – string on tape when machine halts
- Language L is **Turing-recognisable** if some TM **recognises** it, i.e., given an input string x :
 - if x is in L , then the TM halts in state s_Y
 - if x is not in L , then the TM halts in state s_N **or fails to halt**
- Language L is **Turing-decidable** if some TM **decides** it, i.e., given input string x :
 - if x is in L , then the TM halts in state s_Y
 - if x is not in L , then the TM halts in state s_N
- Every decidable language is recognisable, but not every recognisable language is decidable
 - e.g., Halting Problem language
- Enhanced Turing machines
 - 2+ tapes
 - 2D tape
 - non-deterministic
 - (none of these change computing power)
 - Proof: any enhanced TM can be **simulated** by basic TM
- P and NP
 - P is class of DPs solvable by a TM in polynomial time
 - NP is solvable by non-deterministic TM in polynomial time
 - Non-determinism changes only speed of computation

Counter programs

- Completely different model of computation:

- All general purpose programming languages have essentially the same computational power
- Have:
 - variables of type int
 - labelled statements of form:
 - L : unlabelled_statement
 - Unlabelled statements of the form:
 - $x = 0$
 - $x = y + 1$
 - $x = y - 1$
 - if $x == 0$ goto L
 - halt;

Church-Turing Thesis

- The TM is an appropriate model for the 'black box'
 - A whole range of different computational models are equivalent in their abilities

Week 9: Revision

Exam format:

- 4 questions
- 60 marks (15 for each question)
- Duration: 1h 30 mins
 - (extra 30 mins for "uploading")
 - allocate ~20 mins per question
- Examinable material – everything in lectures except:
 - correctness proof for Prim-Jarnik
 - reduction from SAT to Clique
 - proof halting problem is undecidable
 - tutorial questions marked as hard
- Objective: test understanding, not memory
 - **must use own words**
- Advice:
 -
 - if asked to 'describe algorithm X', convince examiner you understand by using a mix of English and pseudocode, possibly with illustrative example
 - Allocation of marks will suggest level of detail expected
 - No own words, no marks
 - Word doc, but draw on paper and insert image/scan into the doc

Problem and problem instances:

- Problem – usually characterised by (unspecified) parameters
- Problem instance – created by giving these parameters values

Class NP

- Decision problems solvable in non-deterministic polynomial time
 - non-deterministic algorithm can make non-deterministic choices
 - apparently more powerful than a normal deterministic algorithm
 - NDA has many possible executions depending on choices made
- NDA:
 - NDA 'solves' a decision problem Π if:
 - for a 'yes' instance, there is **some** execution that returns 'yes'
 - for a 'no' instance, there is **no** execution that returns 'yes'
 - NDA 'solves' in polynomial time if:
 - ... (see earlier lecture)

Polynomial time reductions

- A PTR is a mapping f from a decision problem Π_1 to a DP Π_2 , such that, for every instance i_1 of Π_1 , we have:
 - the instance $f(i_1)$ of Π_2 can be constructed in polynomial time
 - $f(i_1)$ is a 'yes' instance of Π_2 if and only if i_1 is a 'yes' instance of Π_1
- PTR from Π_1 to Π_2 :

$$\Pi_1 \propto \Pi_2$$

NP-complete

- DP Π is NP-complete if:
 - Π is in NP
 - for every problem Π' in NP, Π' is polynomial time reducible to Π
- Consequences of definition
 - If Π is NP-complete and Π is in P, then $P=NP$
 - every problem in NP can be solved in polynomial time by reduction to Π
 - If $P \neq NP$ and Π is NP-complete, then Π is not in P
- Proving:
 - To prove Π_2 is NP-complete, enough to show:
 - Π_2 is in NP
 - there exists a PTR from Π_1 to Π_2
 - How to prove PTR:
 - Figure out which direction the reduction needs to go
 - Suppose we need from Π_1 to Π_2
 - Next, build it:
 - given an instance i_1 of Π_1 , construct an instance of $f(i_1)=i_2$ of Π_2
 - if there is a hint, use it, and otherwise think about how they're related
 - write down/draw instances side by side
 - relate parameters if there are any
 - Finally, show it's correct:
 - Show you can perform construction in polynomial time
 - ... [watch recording] ???

Building automata and Turing machines

- For DFAs, go through each possible option (action) in each state
 - moving to a new state if you have progressed towards the language
 - and moving to an accepting state or sink state if appropriate
- For PDAs, first think about the roles of the stack as a form of memory and then follow a similar approach to DFAs updating the stack
- Defining a TM for even simple problems is hard
 - (won't be asked to make a diagram)
 - Write pseudocode