

## Unix Tutorial 7: Piping Commands

Today, we are going to practise chaining together simple commands using the Unix pipe facility, which is written as a vertical bar i.e. `|`. Most of our pipelines will have the following shape—there is a ‘source’ command that generates some textual output, which is then filtered by a number of other commands. In this tutorial, we will work through two examples.

### Counting Word Occurrences

Download a textfile containing *The Adventures of Sherlock Holmes* from the web, using this command<sup>1</sup>.

```
$ curl -O http://www.gutenberg.org/files/1661/1661-0.txt
$ ls *.txt
```

Notice that a file `1661-0.txt` has been downloaded to your filesystem, into your working directory. You may inspect the contents of file using the `less` command:

```
$ less 1661-0.txt
```

Press space bar to scroll down one screen and `b` to scroll back one screen. Press `q` to quit the `less` program.

Sometimes, you will want to look at just the first (or last) few lines of a long textfile, e.g. a log file. Use the `head` (or `tail`) command for this:

```
$ head -n 10 1661-0.txt
$ tail -n 10 1661-0.txt
```

To print the whole text file out to the console output, use the `cat` command as follows:

```
$ cat 1661-0.txt
```

We are going to use `cat` as the source command in our pipeline. Suppose we want to find all the instances where Holmes says, ‘My dear Watson.’ Let’s build a pipeline that searches for this phrase in our text file:

```
$ cat 1661-0.txt | grep -i "my dear watson"
```

When you run this command, it should print out three lines which are the three occasions when Sherlock utters his memorable phrase to his friend. Notice the `|` is the pipe operator, sending the output from the `cat` command to the `grep` command directly. Notice the `-i` flag for `grep`, which makes the search case-insensitive. There are plenty of other useful flags for `grep`. You might want to

---

<sup>1</sup>Sorry, you will have to type in this command rather than cutting+pastings. The PDF version of this file uses different character encodings for `-` and other characters, that the console will not recognize.

print out a line or two of context before each quotation, or the position in the source file where each expression occurs. Try this command:

```
$ cat 1661-0.txt | grep -inC3 "my dear watson"
```

which will print line numbers for the matching lines, and include 3 lines of context each side of the matching line. Notice that we can run together the flags after a single dash—this saves typing `-i -n -C3` as separate flags.

Here are some more search terms for you—try looking for **elementary** and **cocaine**. How many times does each word occur in the text file?

## Counting Word Occurrences

For a second exercise, we are going to find out the current temperature in Glasgow from the Met Office website. We fetch the Glasgow weather page with a curl command:

```
$ curl https://www.metoffice.gov.uk/weather/forecast/gcuvz3bch
```

This command prints out the whole HTML webpage to the console standard output. We want to filter this command, so it just prints out the relevant information. Let's use the `grep` command to search for the latest observation of weather<sup>2</sup>.

```
$ curl -s https://www.metoffice.gov.uk/
weather/forecast/gcuvz3bch | grep tempText |
head -n 1
```

Note the `-s` flag to curl makes it operate *silently* so it does not print out its progress information as it downloads the data. The `grep` command prints out matching lines for the `tempText` paragraph class.

Now we need to use some search-and-replace commands (`sed`) to strip out the HTML tags, with their angle-bracket start and end characters:

```
$ curl -s https://www.metoffice.gov.uk/
weather/forecast/gcuvz3bch | grep tempText |
head -n 1 | sed -e 's/<[^>]*>/ /g'
```

This returns the headline weather forecast for Glasgow.

---

<sup>2</sup>This is a long command, so it may spread over multiple lines, but you can carry on typing directly on a single line that gets wrapped by the console.

Finally, let's save the weather report to a textfile called `glasgowtemp.txt`. We can use the `>` operator to redirect the output of the pipeline to a file.

```
curl -s https://www.metoffice.gov.uk/  
weather/forecast/gcuvz3bch | grep tempText  
| head -n 1 | sed -e 's/<[^>]*>/ /g' >  
glasgowtemp.txt
```

Some more challenges:

1. (various approaches possible) Adapt this pipeline to print out the temperature in London.

## Further Reading

For further investigation today, use the `man` command to find out the flags for all the commands you have already executed.

There are lots of `grep` tutorials online, e.g. see <http://www.thegeekstuff.com/2009/03/15-practical-unix-grep-command-examples/>.