



INDEXING METHODOLOGY: II

DISCUSSION WEEK 7

Database Systems (H)
Dr Chris Anagnostopoulos



B+ TREE ON NON-ORDERING NON-KEY!

- $n = 1000$ departments; $r = 100,000$ employees, e.g., BT Group
- Blocking factor **bfr** = 50 records/block
- DNO is *uniformly* distributed: $d = r/n = 100$ employees/department
- Leaf node order: $q = 10$ DNO values/data-pointers
- Tree node order: $p = 10$ tree-pointers
- 100 pointers can fit in 1 block.

Task 1: How many leaf nodes do we need for *all* DNO values?

Task 2: Which is the B+ Tree *leaf* structure, the *whole* B+ Tree structure and finally the *expected cost* searching for the employees of Dept. 3.



B+ TREE ON NON-ORDERING NON-KEY!

- $n = 1000$ departments; $r = 100,000$ tuples (employees)
- Blocking factor **bfr = 50 records/block**
- DNO is *uniformly* distributed: $d = r/n = 100$ employees/department
- Leaf node order: $q = 10$ DNO values/data-pointers
- Tree node order: $p = 10$ tree-pointers
- 100 pointers can fit in 1 block.

The file has: $b = \text{ceil}(100,000/50) = 2000$ blocks

Task 1: How many leaf nodes do we need for *all* DNO values?


- $n = 1000$ departments (DNOs) can fit in $n/q = 100$ leaf nodes; each leaf node has $q = 10$ depts.

Task 2: We have 100 leaf nodes; each one with 10 depts.



Leaf Node Structure:

- For *each* DNO value, assign a pointer to *block(s)* with *block of pointers* of those blocks with employees working at this DNO (dept.)
- We need $d = 100$ block pointers per DNO (dept.) stored in 1 block
- We need $m = 10$ **blocks** of block of pointers (one block per DNO value), each block with 100 block pointers (100 pointers fit in one block)



1 Leaf node (10 depts.)

10 blocks of block-pointers

100 data-blocks/dept.

data-blocks for employees in DNO=3
(100 employees work in dept. 3 residing in 100 different
blocks; worst-case)

Task 2: Root: 1 node ($p = 10$); L1: 10 nodes; Lf = 100 nodes (level $t = 3$)



Root: 1 block



...




L1: 10 blocks




...

...

Lf: 100 blocks



UG: $100 \times 10 = 1000$ blocks



Index Size: 1111 blocks.

File Size: 2000 blocks.

Index I/O: $1 + 1 + 1 + 1 + 100 = 104$ block accesses

File I/O: 2000 block accesses

94.8% speed-up

55% storage